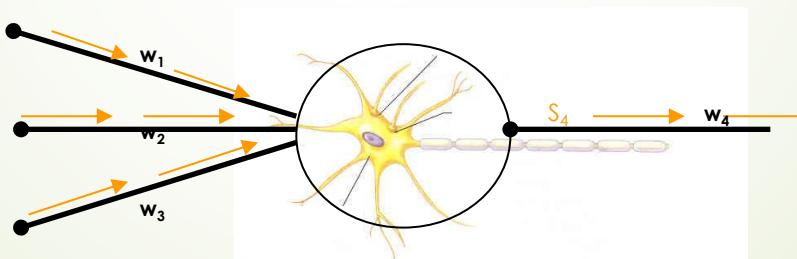


Artificial Neural Networks, ANN

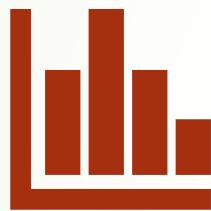
Mikael Pontarp
Lund University



How ANNs Are Used in Biology

Inference vs.
Mechanistic
Modeling

Two Fundamental Uses of ANNs



1. Statistical Inference
(Most Common)



2. Forward Mechanistic
Modeling (Less Common)

ANNs as Statistical Inference Tools

Used to detect patterns, classify samples, and predict outcomes:

- ▶ Crow hybrid classification
 - ▶ E.g. identify and classify individuals
- ▶ Cancer prognosis prediction
 - ▶ E.g. Identify tumor subtypes
- ▶ Protein structure prediction (AlphaFold)
 - ▶ E.g. infer 3D structure from amino acid sequence
- ▶ Gene expression prediction
 - ▶ E.g. predict expression based on promoter sequences



Hybridization in the crow

- ▶ Carrion crow and hooded crow hybridize across Europe.
- ▶ Hybrid zone produces intermediate phenotypes.
- ▶ Classification task: determine species from image data.
- ▶ ANN needed because classification relies on pixel patterns.



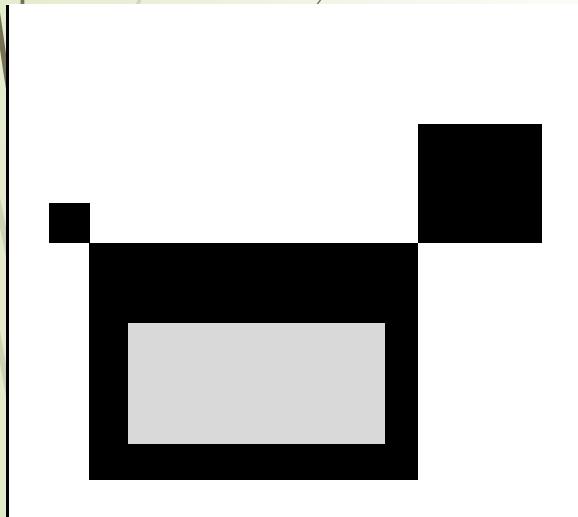
Input signals from "pixel" images

$$S_1 = -1 = \boxed{\text{light gray}}$$

$$S_2 = 0 = \boxed{\text{white}}$$

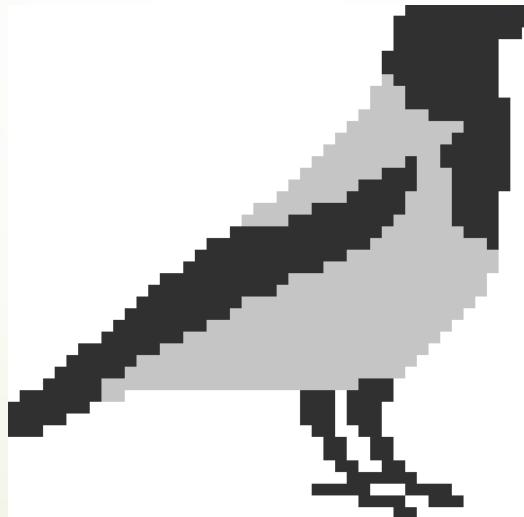
$$S_3 = 1 = \boxed{\text{black}}$$

14 x 14



196

45 x 45



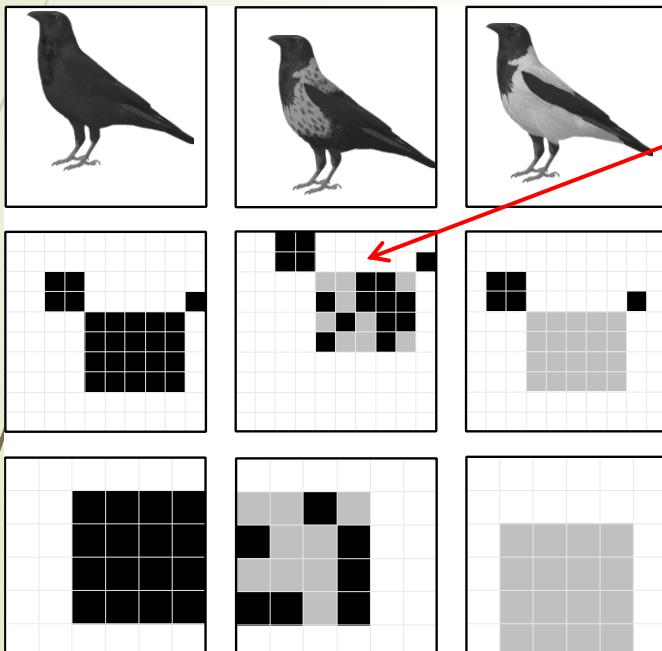
2025

200 x 200



40000

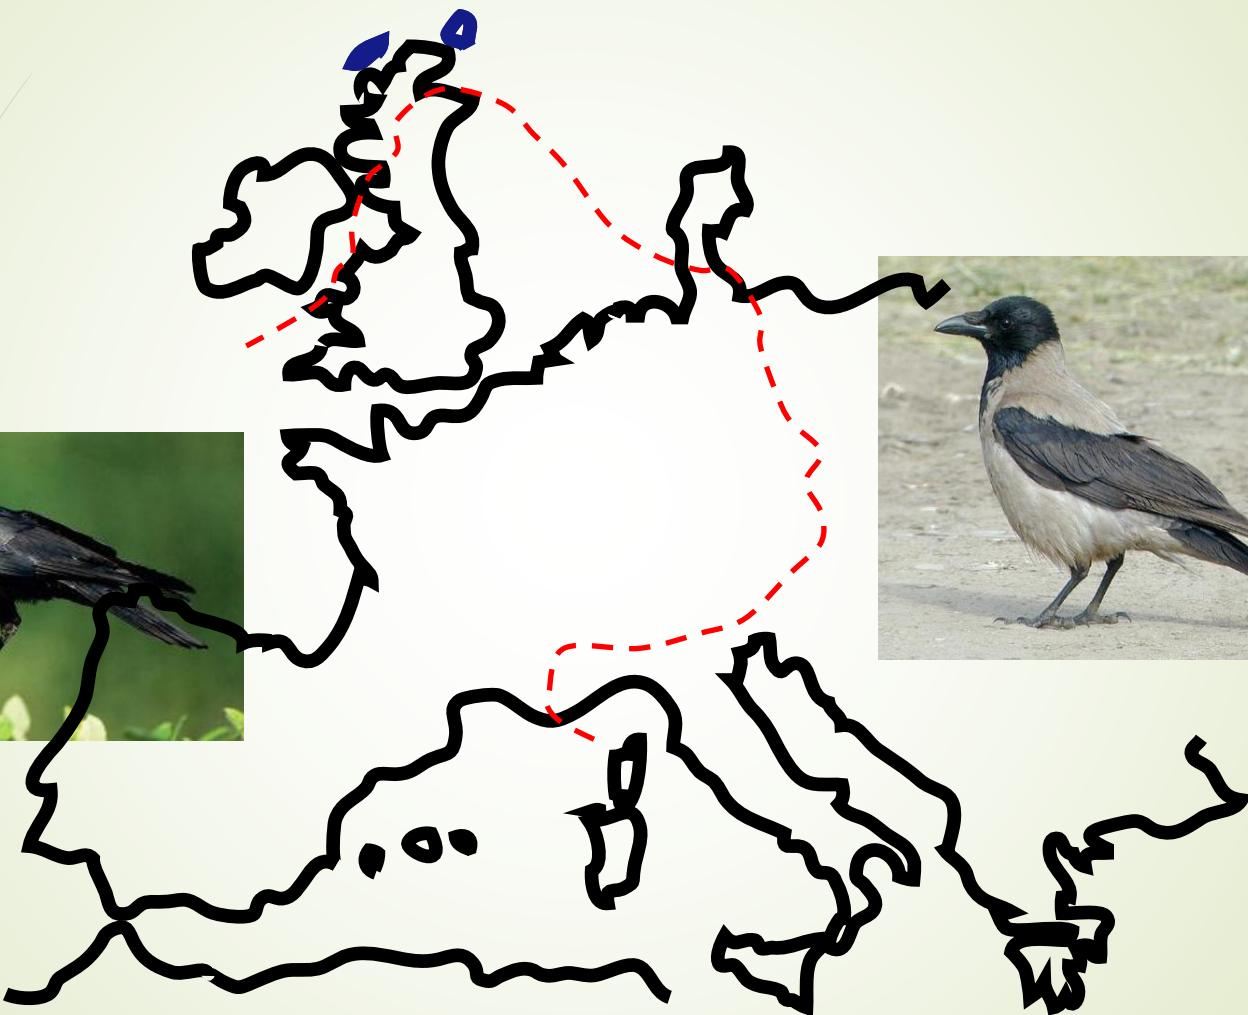
What information is the network learning?

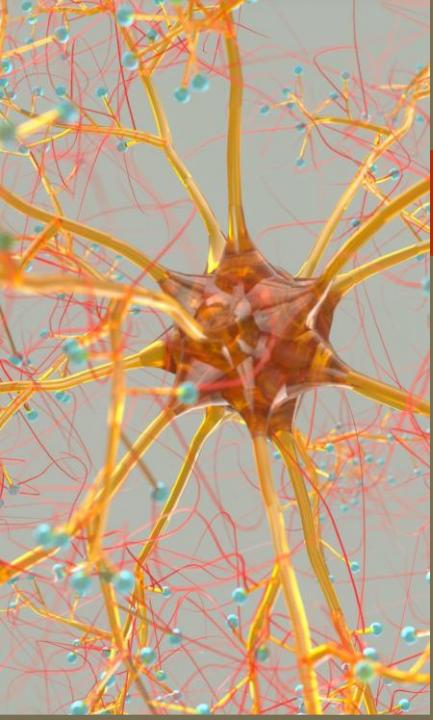


0011000000011000010000-1 -1 ...
000000000001100000001100...

Images as ANN Inputs: Pixels Become Signals

- ▶ Images are converted into numerical input values.
- ▶ Each pixel → a signal (e.g., black=1, gray=0, white=-1).
- ▶ Higher resolution → more input neurons:
 - ▶ – $14 \times 14 \rightarrow 196$ pixels
 - ▶ – $45 \times 45 \rightarrow 2,025$ pixels
 - ▶ – $200 \times 200 \rightarrow 40,000$ pixels
- ▶ ANN must learn patterns in these high-dimensional vectors.

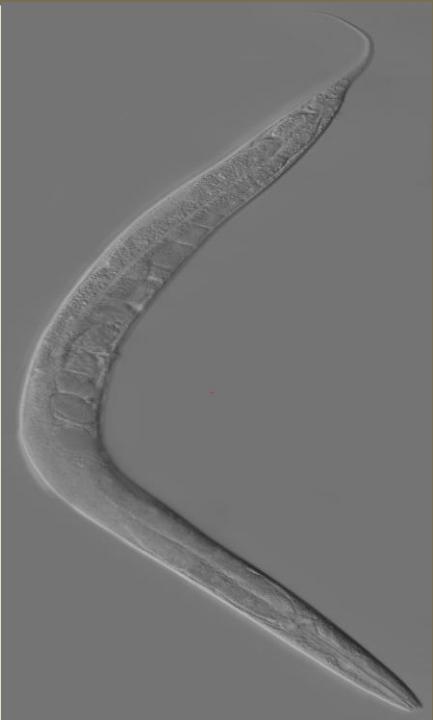




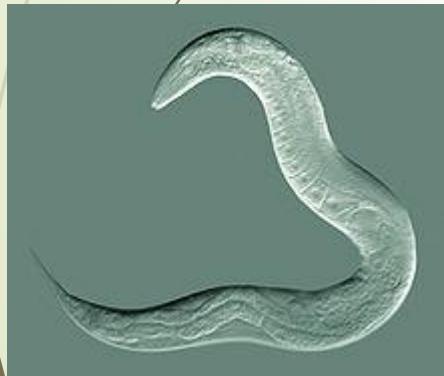
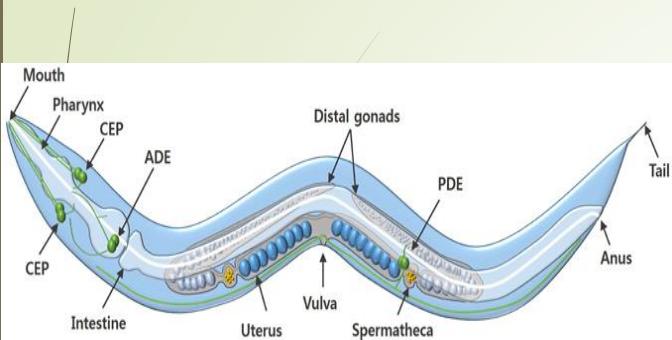
ANNs as Forward Models

Used to simulate biological processes or decisions:

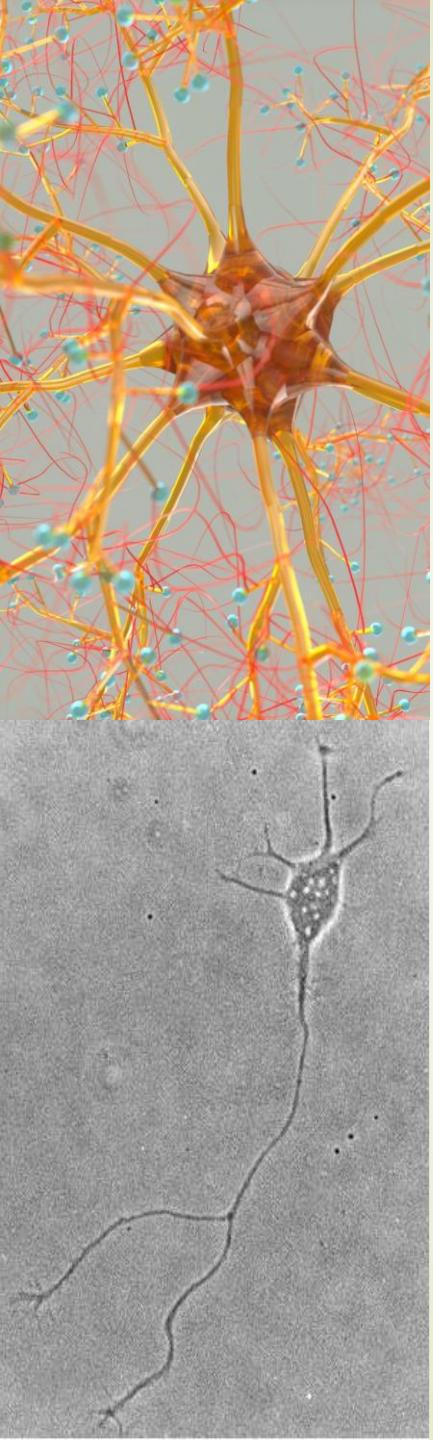
- ▶ Neuronal logic (AND/OR/XOR)
 - ▶ E.g. gene regulation, immune response, sex determination
- ▶ Simulate movement, sensory response, etc
 - ▶ E.g. *C. elegans* neural circuit
- ▶ Retinal feature detection models
 - ▶ E.g. edges, motion, light
- ▶ Models of animal behavior
 - ▶ E.g. navigation, foraging, defense systems



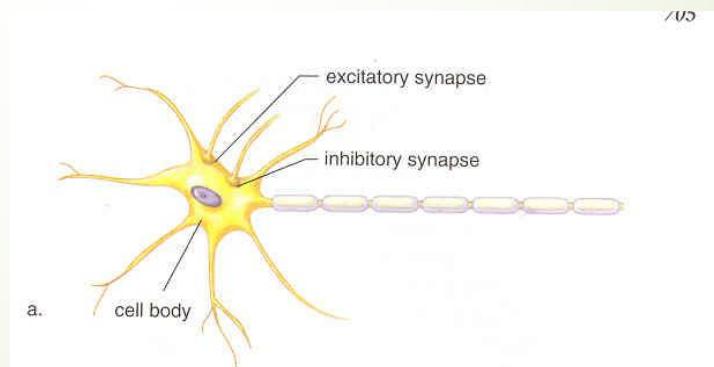
ANNs as Forward Models

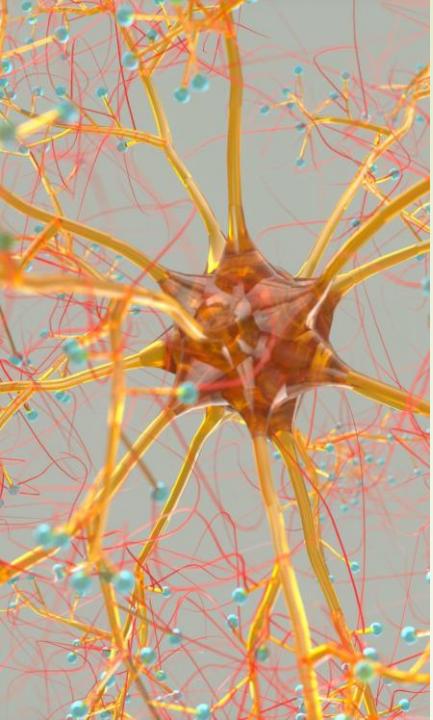


- ▶ *C. elegans* has only 302 neurons — a fully mapped biological neural network.
- ▶ Its connectome inspires ANN models that simulate behavior from neural wiring.
- ▶ ANNs can mimic sensory input → internal processing → motor output.
- ▶ Forward simulation: given stimuli, the ANN predicts organism responses.
- ▶ Enables testing hypotheses about behavior, decision-making, and neural control.
- ▶ Demonstrates how artificial networks can approximate real biological systems.

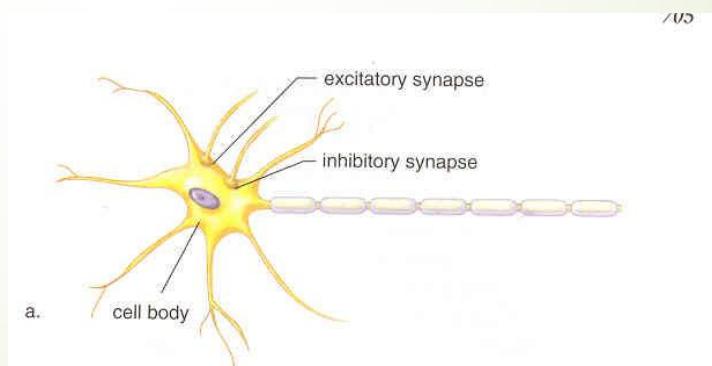
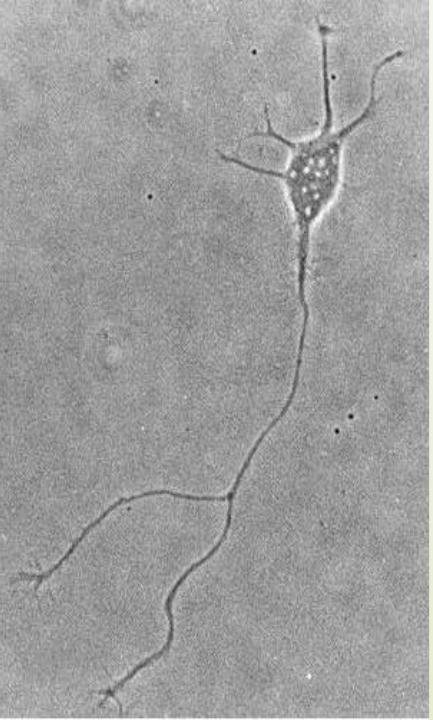


- ▶ ANN architecture is biologically inspired
- ▶ ANN usage is usually statistical
- ▶ ANN interpretation can be mechanistic
- ▶ Sits between ML, systems biology, and neuroscience





- a neuron accumulates signals from synapses on dendrites and soma
- fires at some threshold
- spike that propagates to other neurons in axon



How Simple Neurons Create Complex Decisions

- ▶ A single artificial neuron can only draw one straight decision boundary.
- ▶ It can learn simple rules like AND or OR, but not XOR.
- ▶ When neurons are connected in layers, each learns a simple rule (a simple boundary).
- ▶ By combining these simple boundaries, the network forms complex, nonlinear decisions.
- ▶ This is why multilayer networks can solve XOR and other nonlinear problems—many simple neurons working together create powerful behavior.

Summary: AND, OR, and XOR Logic

- ▶ AND (“and”): Output = 1 only when BOTH inputs are 1.
 - Linearly separable → a single neuron can learn it.
- ▶ OR (“or”): Output = 1 when AT LEAST ONE input is 1.
 - Also linearly separable → a single neuron can learn it.
- ▶ XOR (“exclusive or”): Output = 1 when EXACTLY ONE input is 1.
 - NOT linearly separable → requires multiple neurons (hidden layer).

Biological Examples of AND, OR, and XOR Logic

- ▶ AND: Gene expression requires two transcription factors
 - Example: A gene activates only when TF1 AND TF2 bind.
- ▶ OR: Immune cells activate if any danger signal is detected
 - Example: viral RNA OR fungal β -glucan triggers inflammation.
- ▶ XOR: Output depends on exactly one of two signals
 - Example: Reptiles sex determination, extreme (high or low) temperature \rightarrow female while intermediate temperature \rightarrow male.

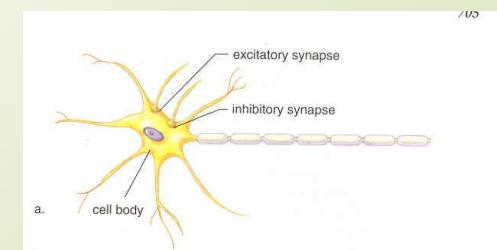
Artificial neuron:

dendrites

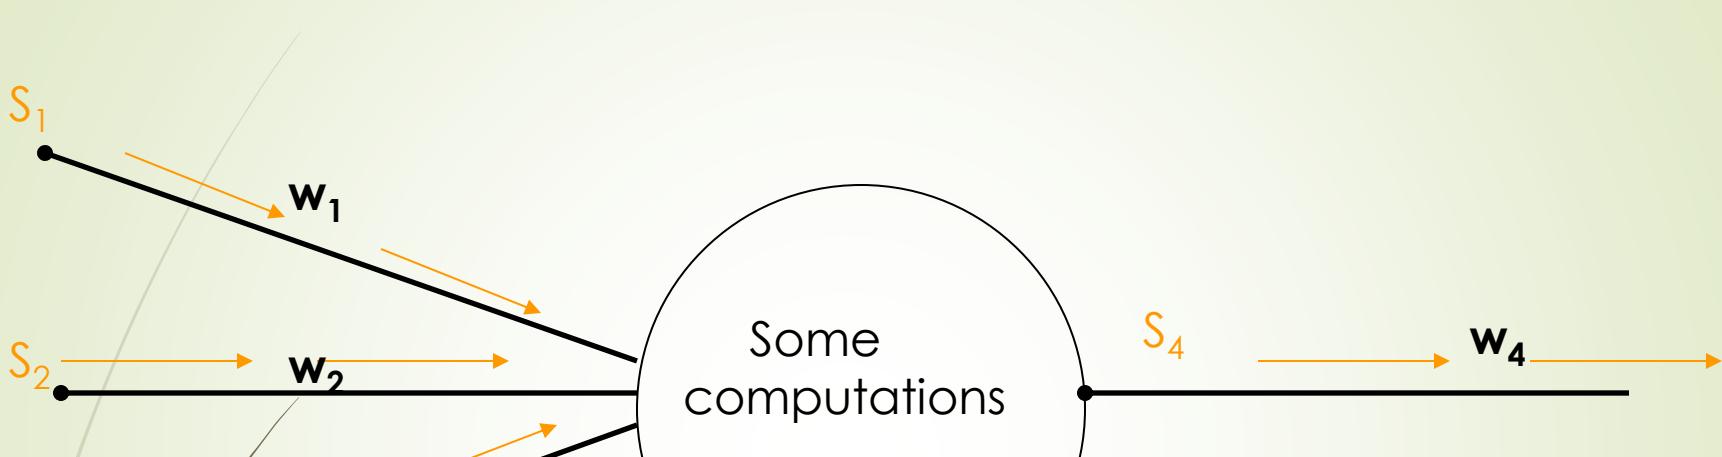
soma

axon

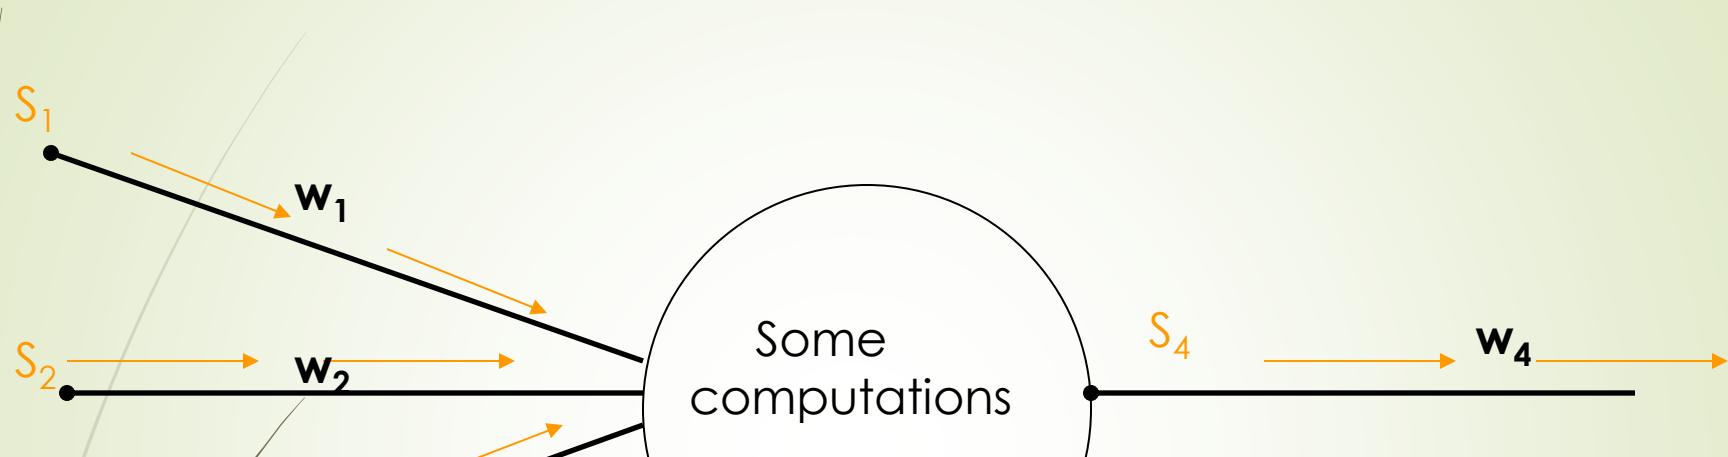
dendrites = axons!



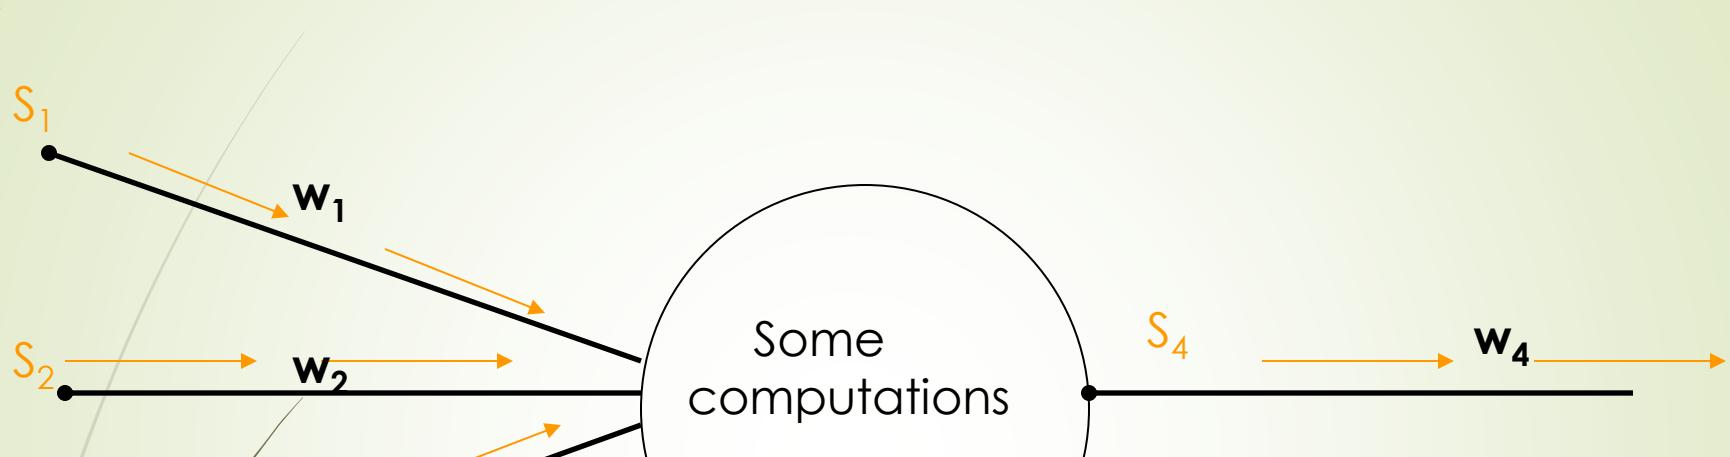
Dendrites and axons in real neurons, weights (w) in artificial ones:



Input: Three input signals, S_1, S_2, S_3



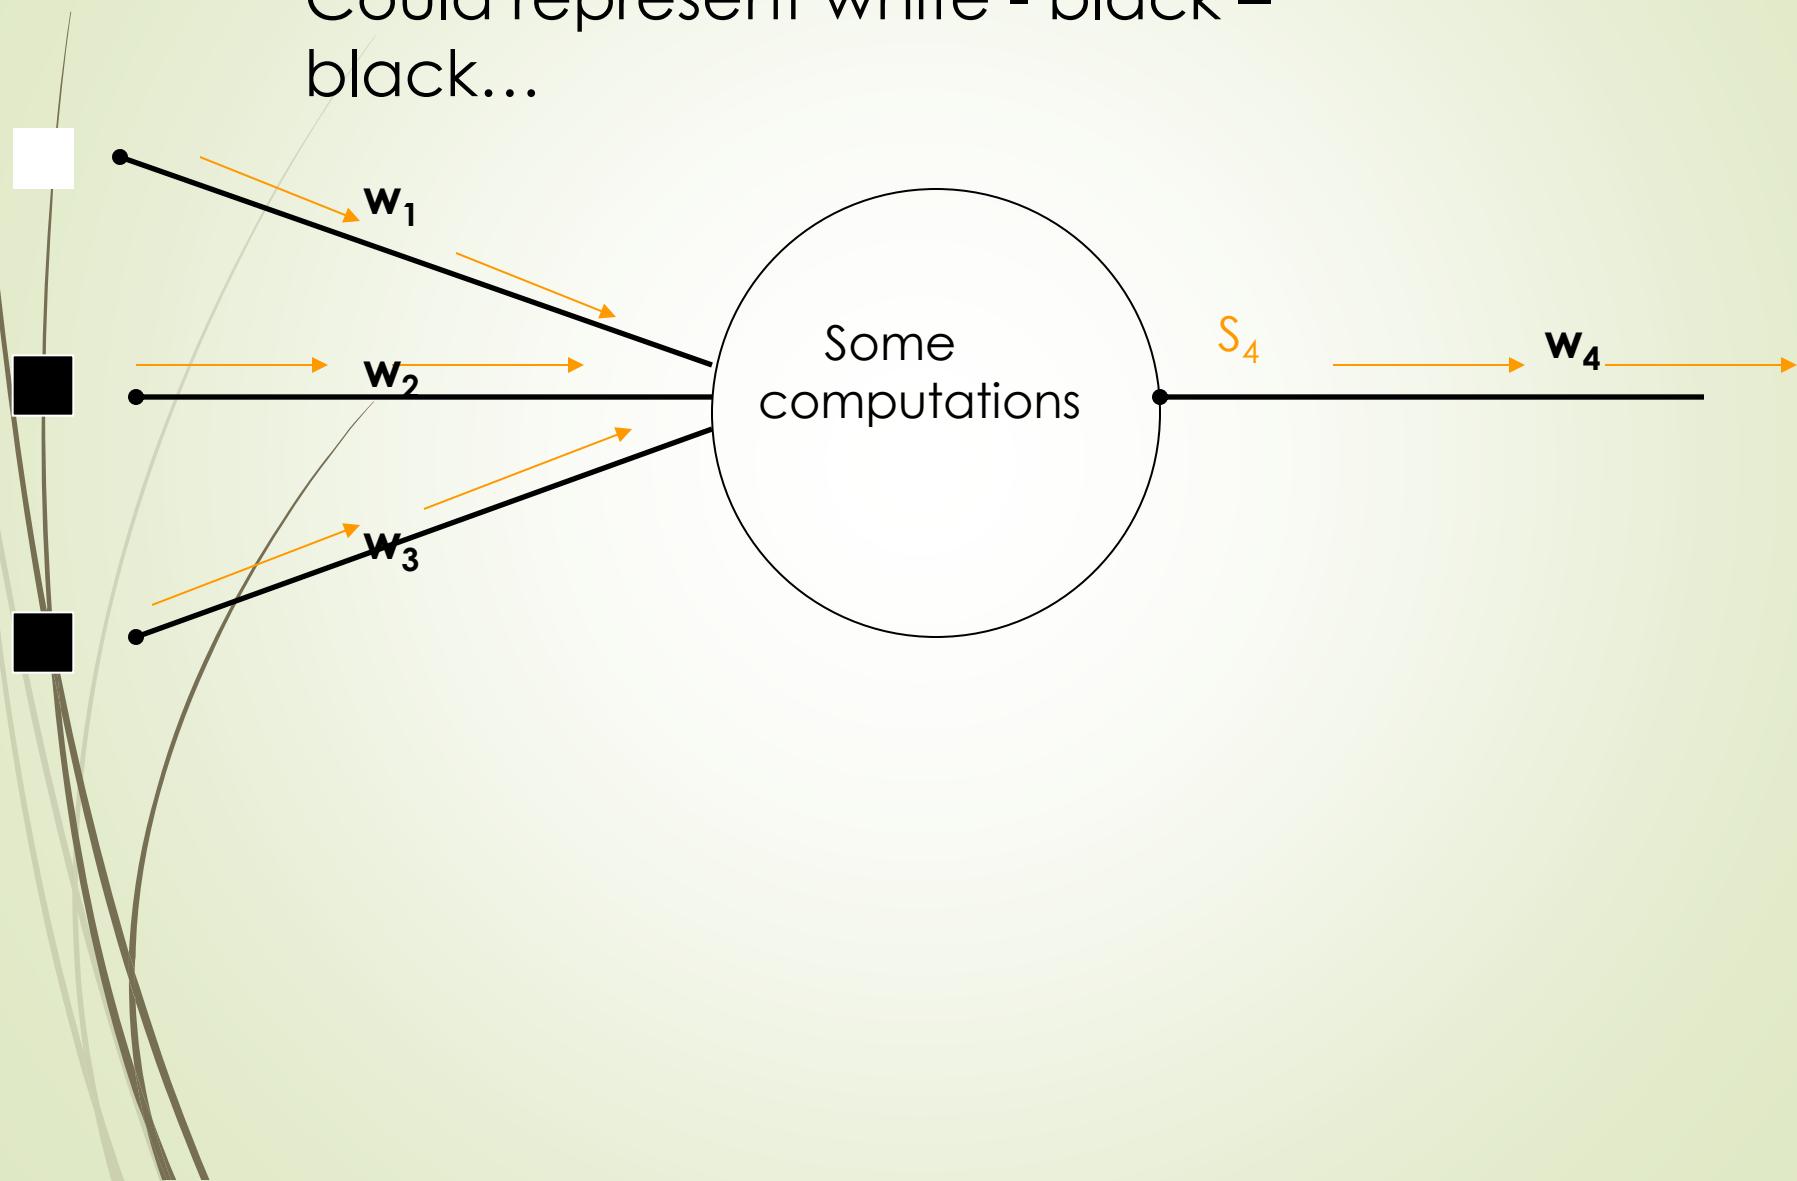
Input: Three input signals, S_1, S_2, S_3



Output: one signal, S_4

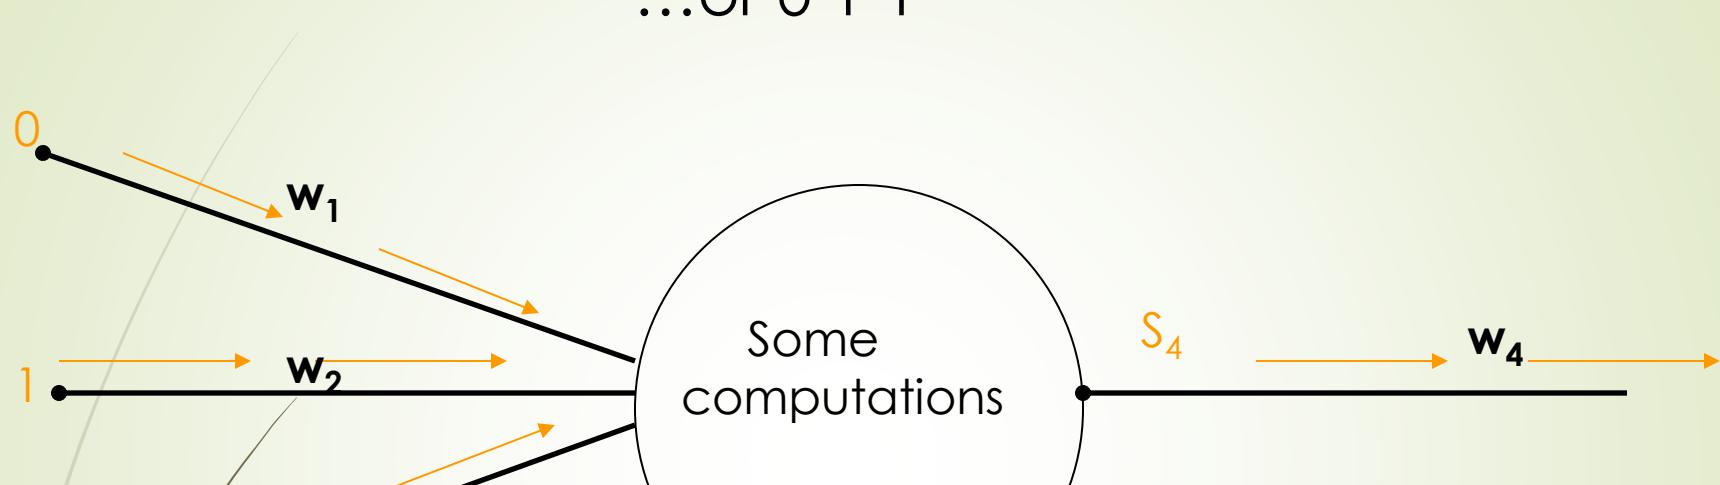
Input: Three input signals, S_1, S_2, S_3

Could represent white - black – black...



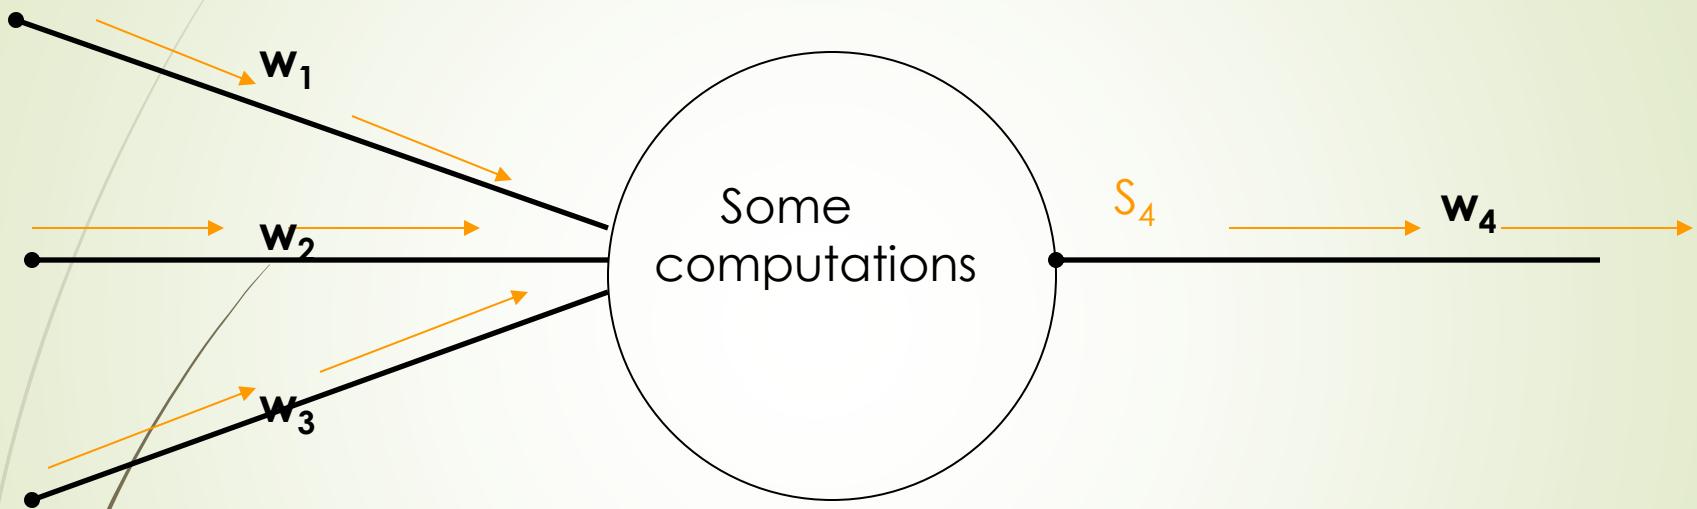
Input: Three input signals, S_1, S_2, S_3

...or 0 1 1

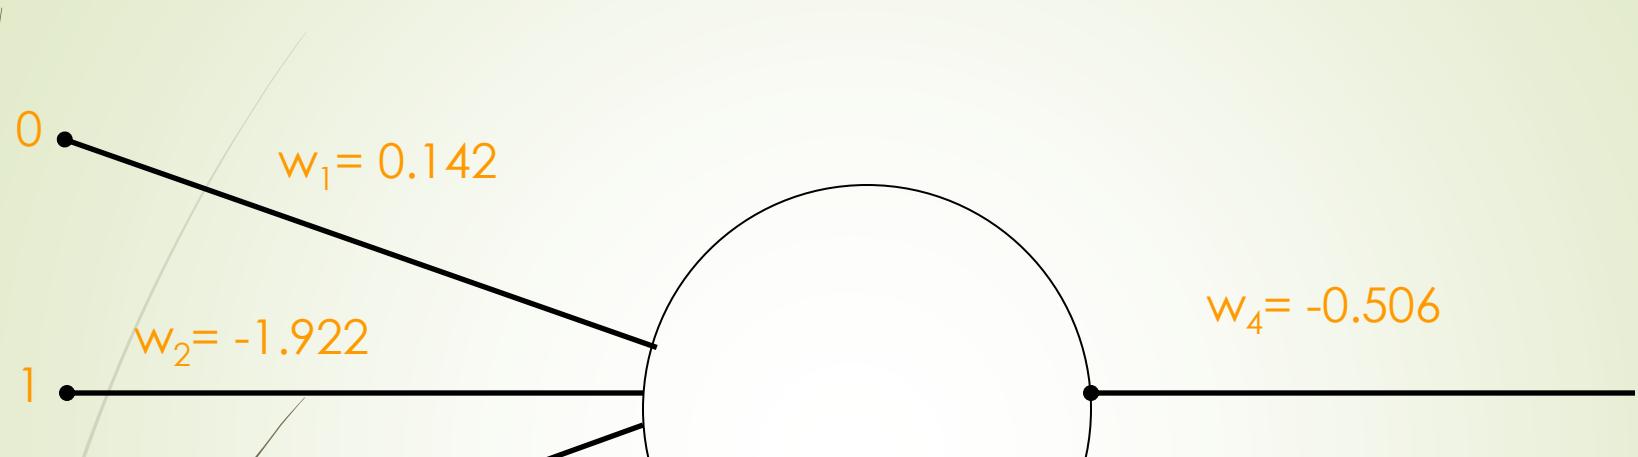


Input: Three input signals, S_1, S_2, S_3

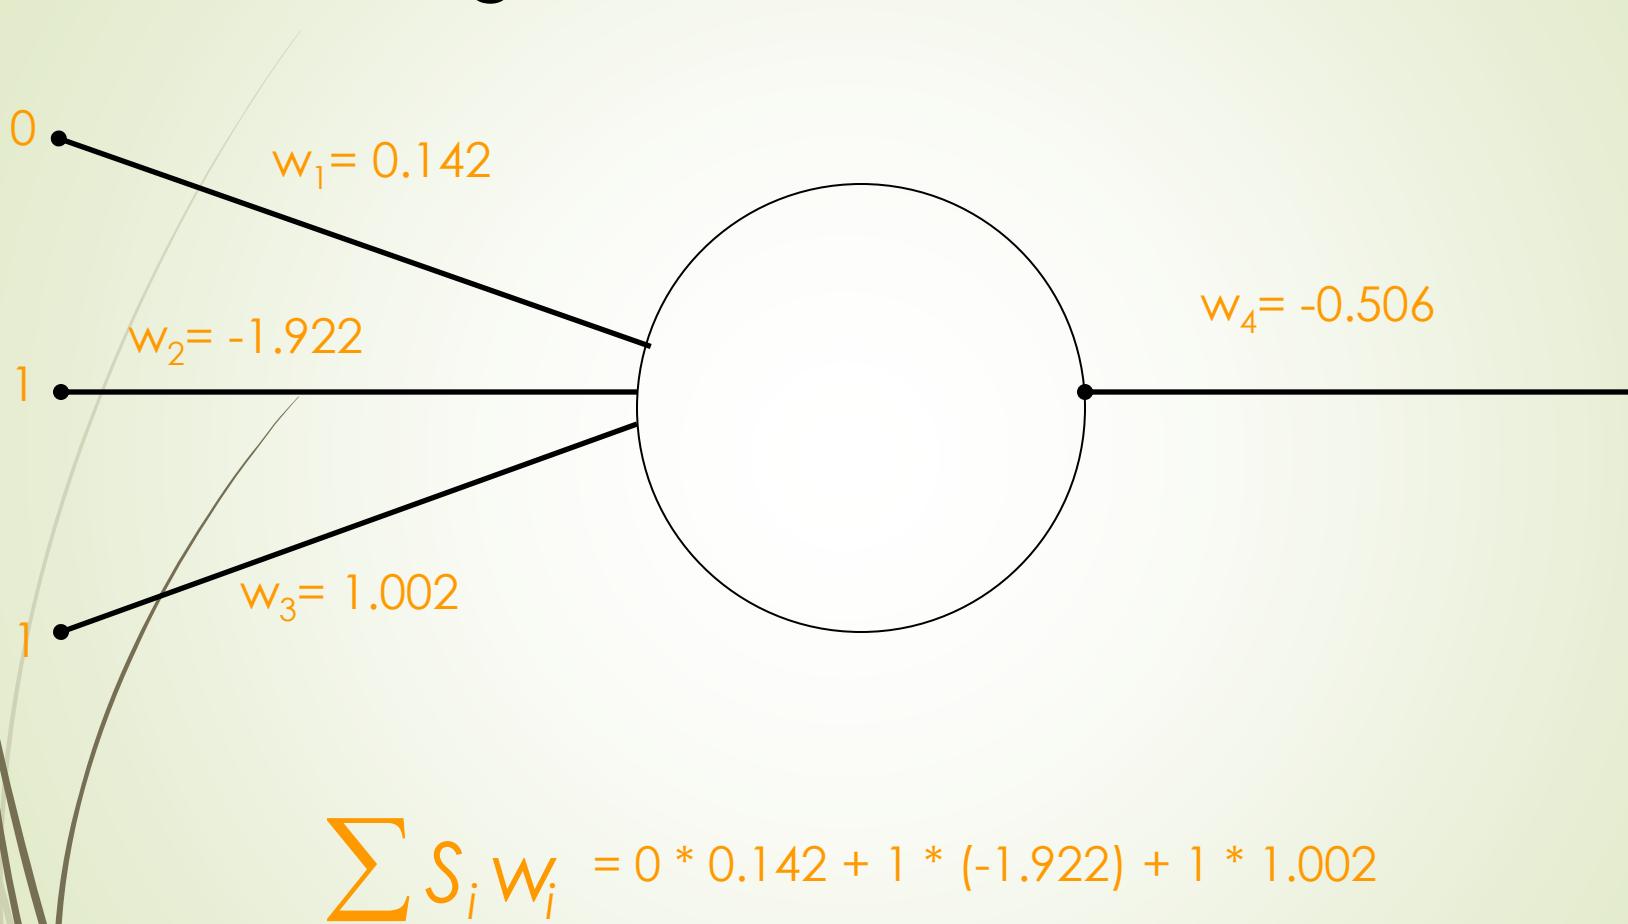
...or more often continuous values



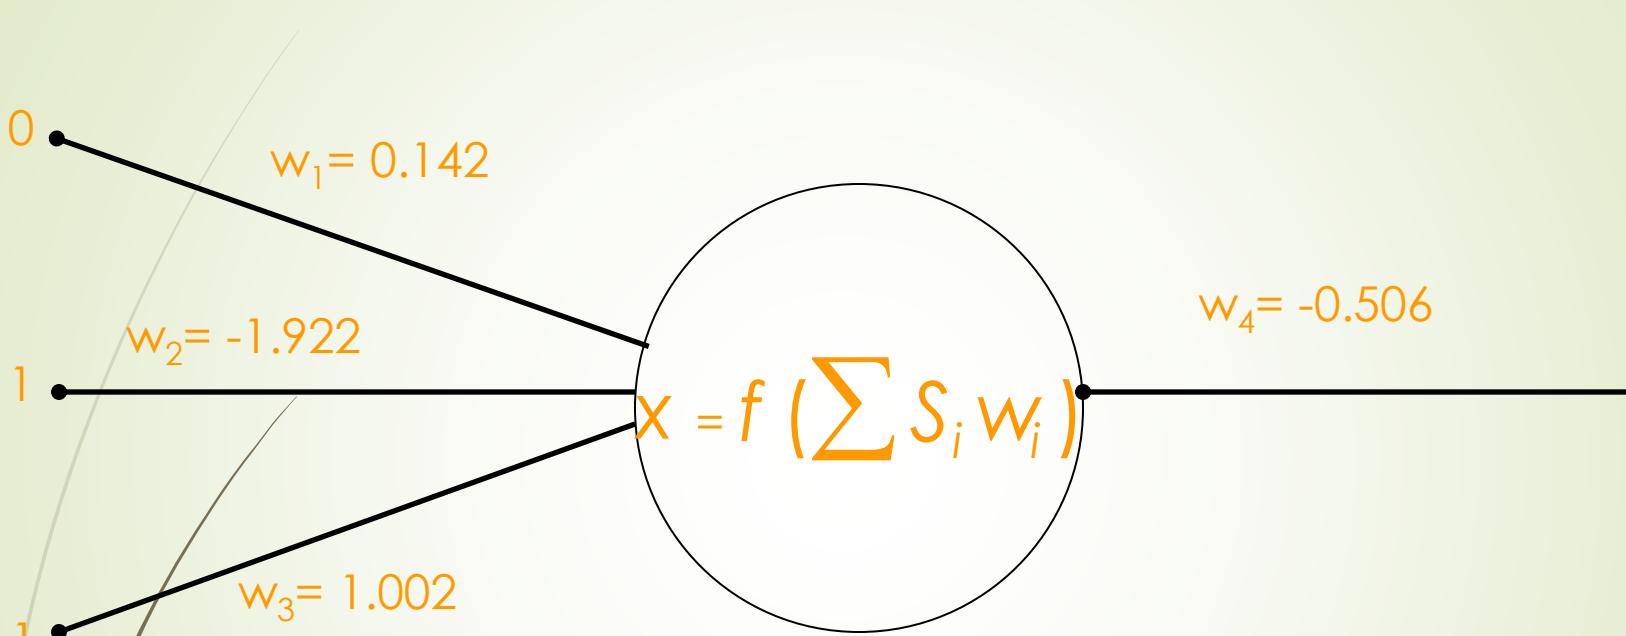
Set weights to random decimal numbers



Multiply with input signals

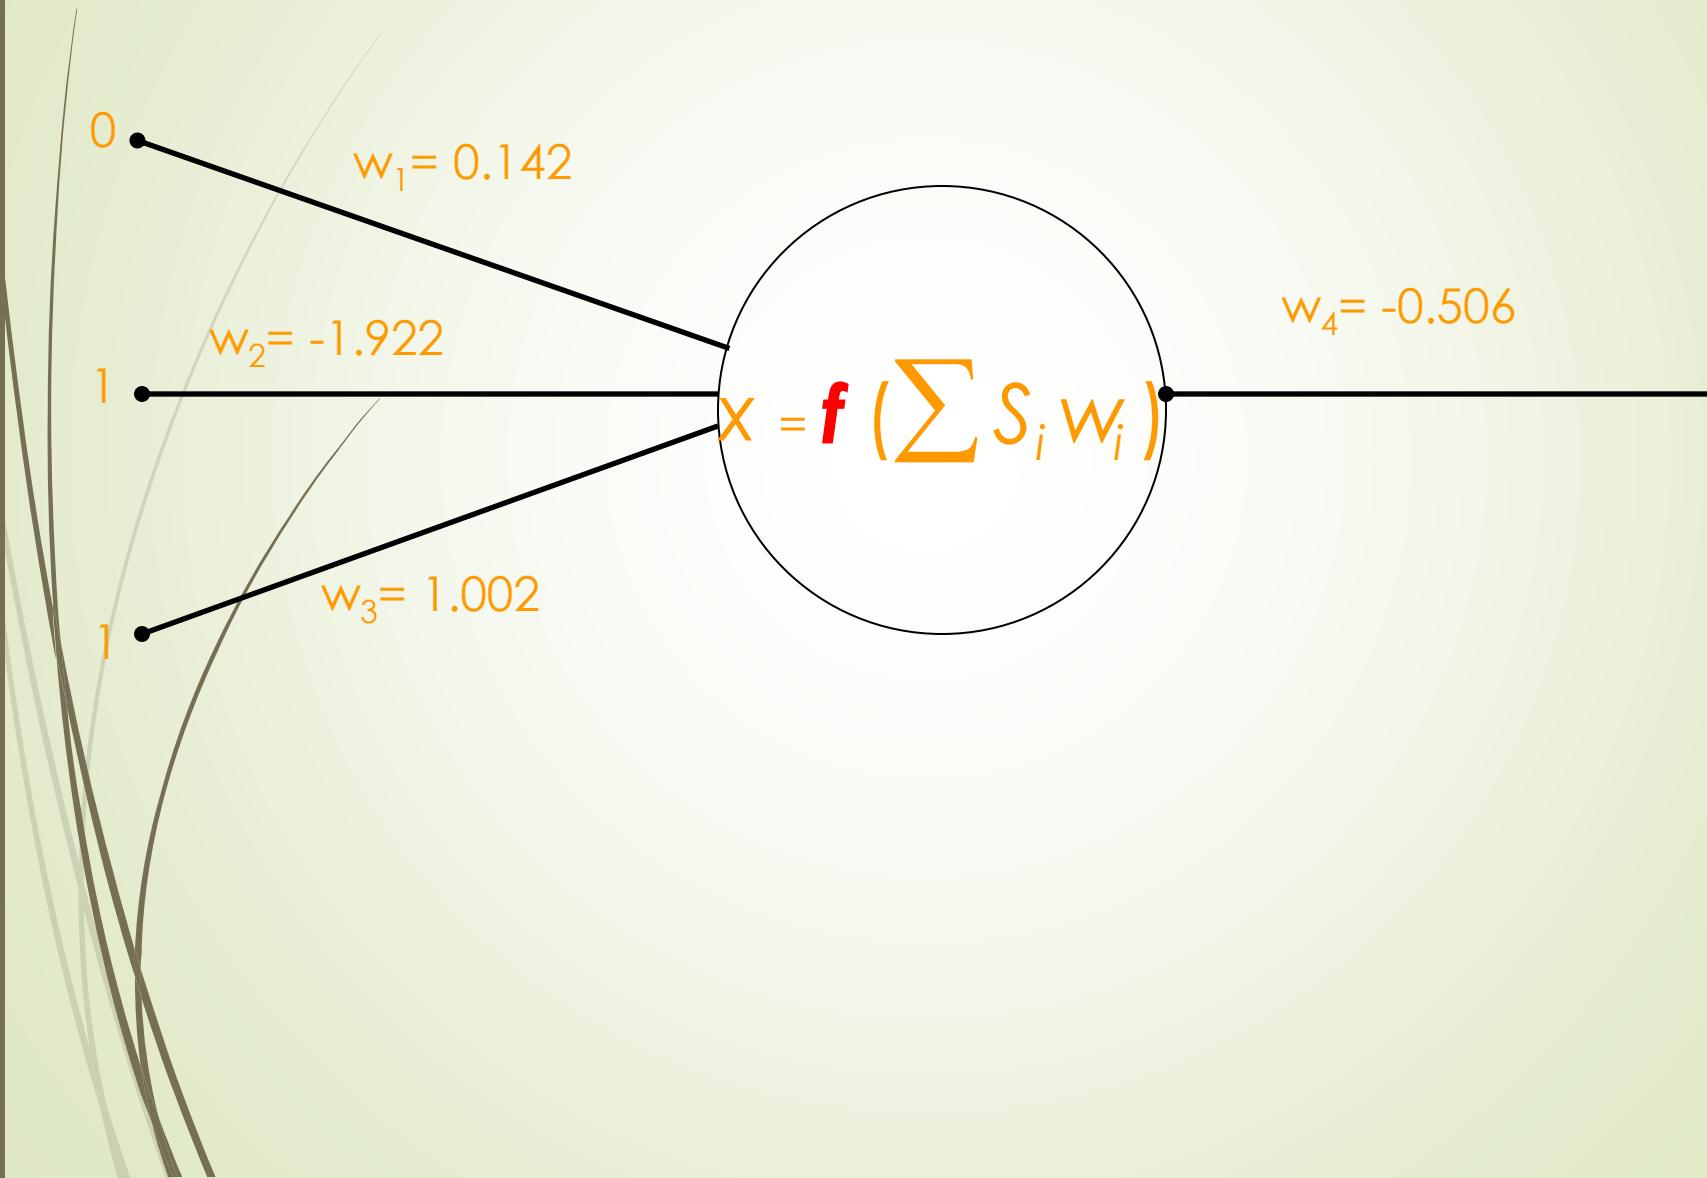


This happens in the "soma" =
node = neuron body



$$\sum s_i w_i = 0 * 0.142 + 1 * (-1.922) + 1 * 1.002$$

Activation function decides how signal propagates



Activation (or transfer) functions

1

0

-5

0

5

$$f(T) = \begin{cases} 0 & \text{if } T \leq 0 \\ 1 & \text{if } T > 0 \end{cases}$$

1

0

$$f(T) = \frac{1}{1 + e^{-T}}$$

1

0

-1

$$f(T) = \tanh(x)$$

Activation (or transfer) functions

1

0

-5

0

5

1

0

1

0

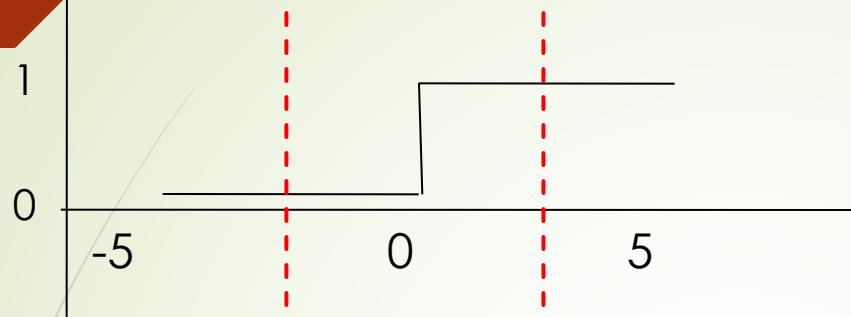
-1

$$f(T) = \begin{cases} 0 & \text{if } T \leq 0 \\ 1 & \text{if } T > 0 \end{cases}$$

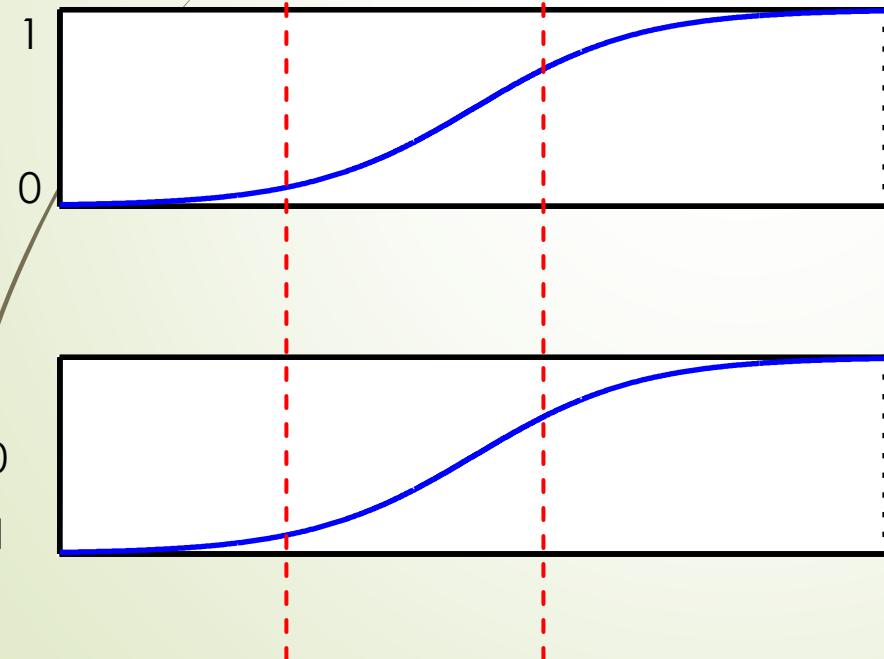
$$f(T) = \frac{1}{1 + e^{-T}}$$

$$f(T) = \tanh(x)$$

Activation (or transfer) functions



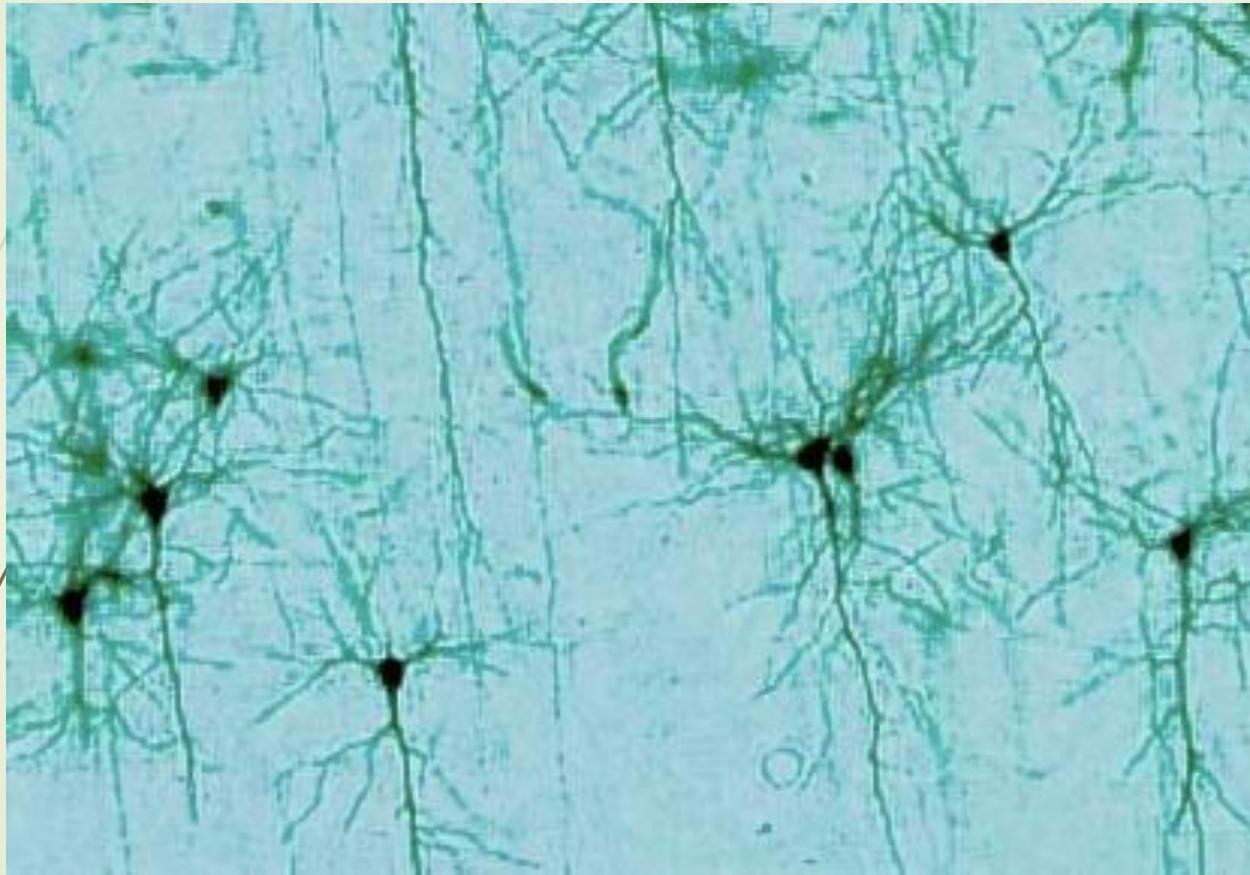
$$f(T) = \begin{cases} 0 & \text{if } T \leq 0 \\ 1 & \text{if } T > 0 \end{cases}$$



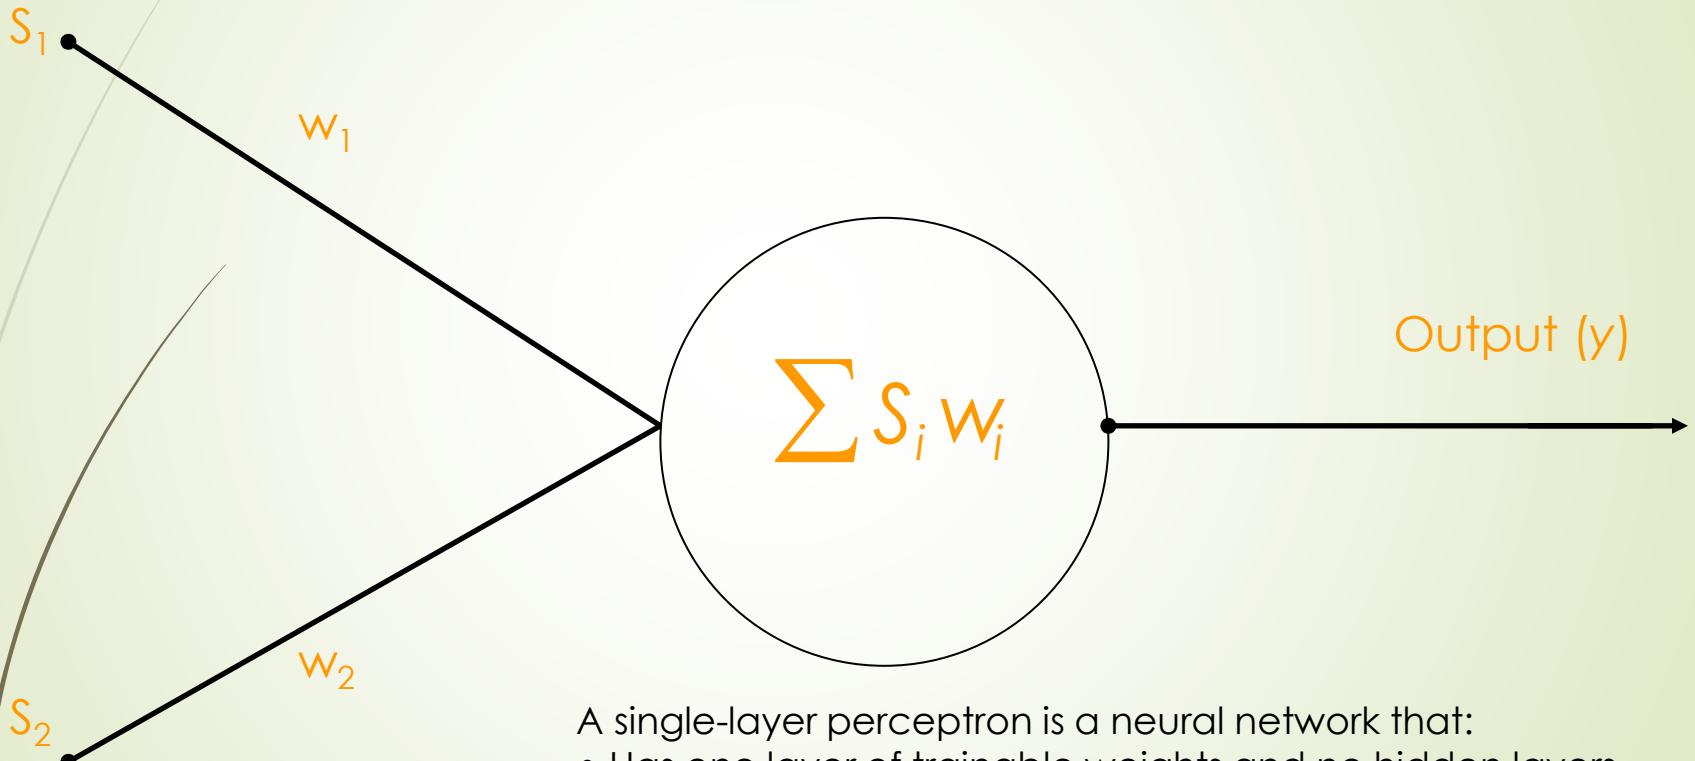
$$f(T) = \frac{1}{1 + e^{-T}}$$

$$f(T) = \tanh(x)$$

Real network:



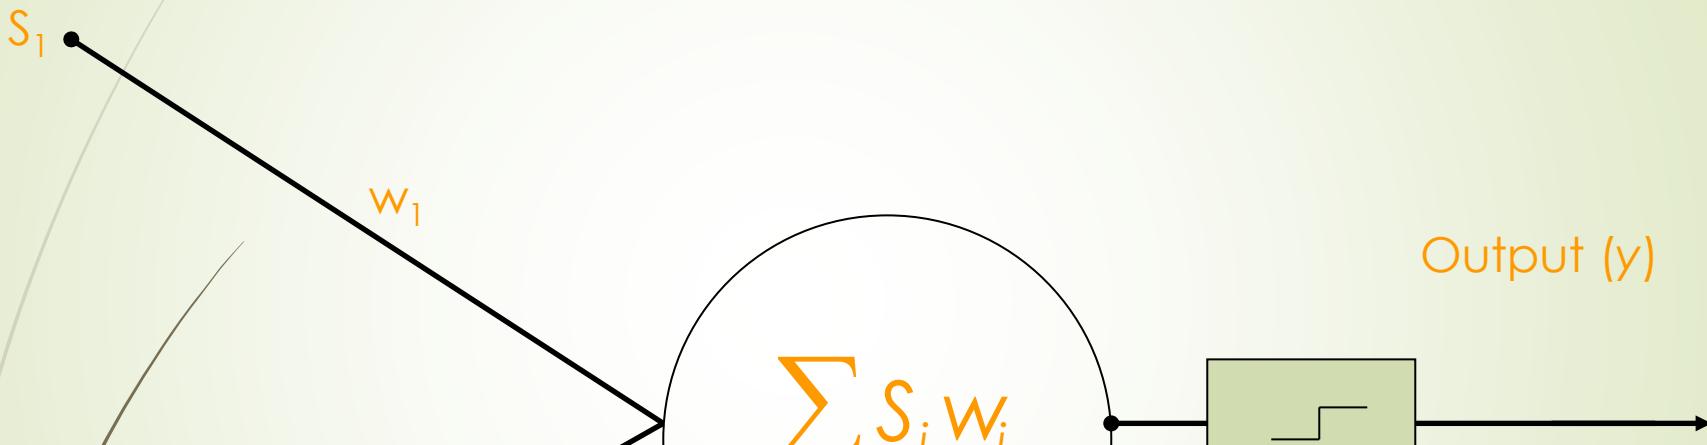
The simplest network: a single-layer perceptron



A single-layer perceptron is a neural network that:

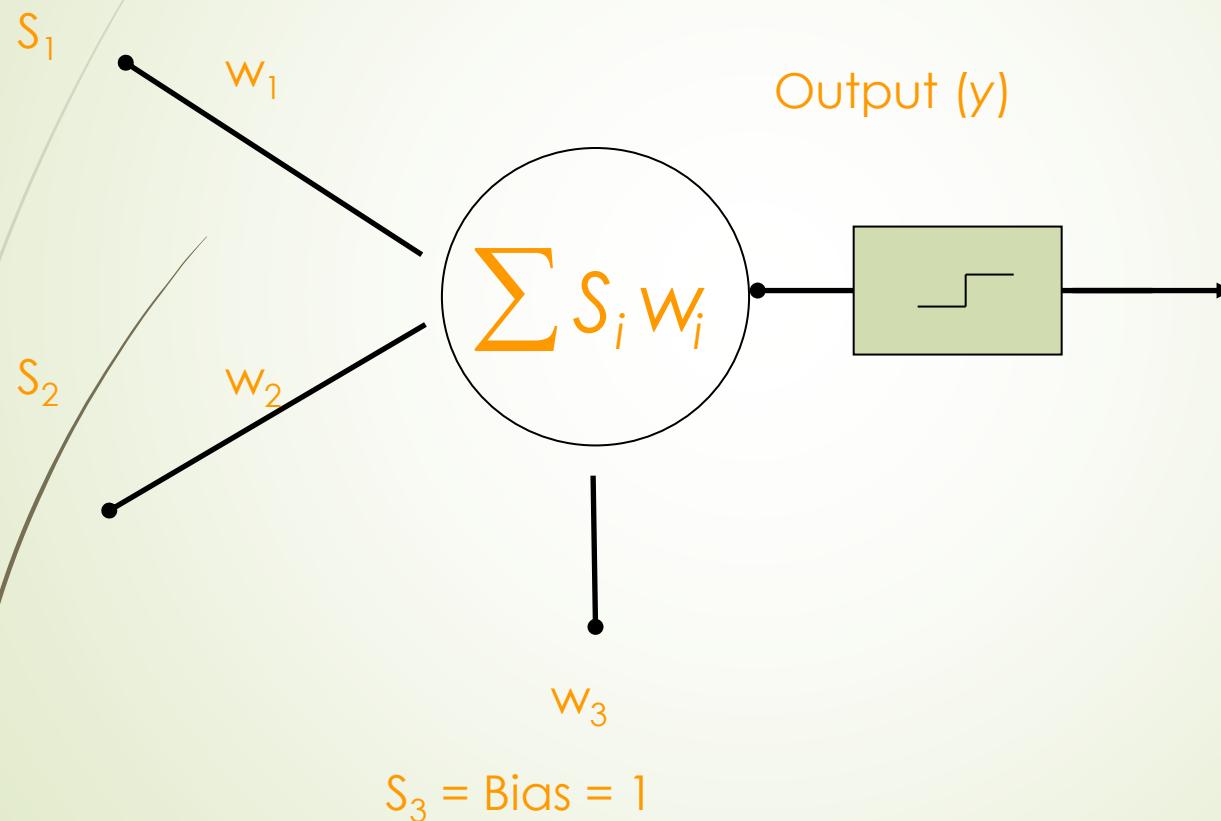
- Has one layer of trainable weights and no hidden layers
- Contains one or more neurons receiving all input features directly
- Computes a weighted sum and applies an activation function
- Can represent only linear decision boundaries

Activation function



$$y = \begin{cases} 1 & \text{if } \sum s_i w_i + b > 0 \\ 0 & \text{else} \end{cases}$$

Bias and truth table



s_1	s_2	out
0	0	?
1	0	?
0	1	?
1	1	?

Bias in a Perceptron

- ▶ The bias acts like a constant input always equal to 1
- ▶ It allows shifting the decision boundary left or right
- ▶ Equivalent to a neuron's threshold
- ▶ Without bias, many functions cannot be learned

Truth Table Use in Training

- ▶ The truth table lists all possible inputs and desired outputs
- ▶ For each row, compute: $T = \sum(S_i w_i) + b$
- ▶ Apply activation function → prediction
- ▶ If prediction ≠ target → update weights
- ▶ Example (AND):
- ▶ $(0,0 \rightarrow 0), (1,0 \rightarrow 0), (0,1 \rightarrow 0), (1,1 \rightarrow 1)$



Training it on logical operators:

AND:

s_1	s_2	out
0	0	0
1	0	0
0	1	0
1	1	1

OR:

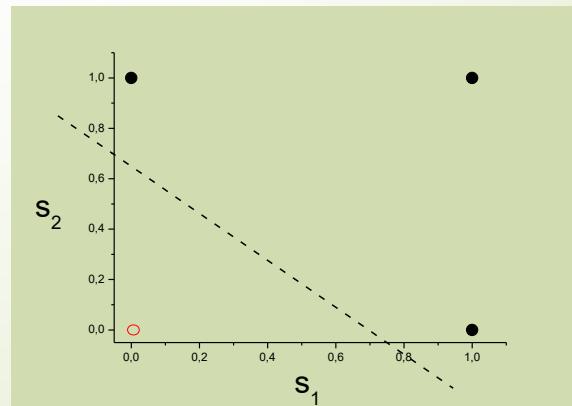
s_1	s_2	out
0	0	0
1	0	1
0	1	1
1	1	1

XOR:

s_1	s_2	out
0	0	0
1	0	1
0	1	1
1	1	0

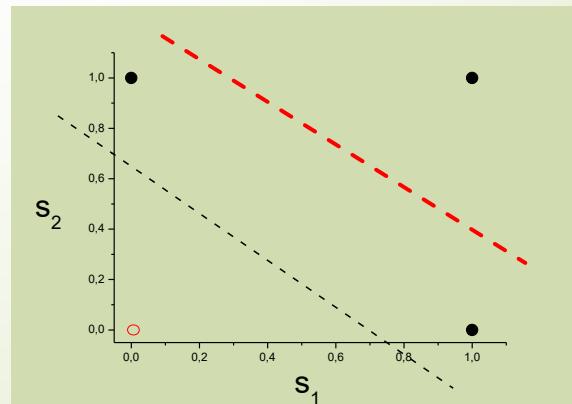
Graphical Representation of the OR Function

- Inputs s_1 and s_2 plotted in 2-D plane
- Black dots = $\text{OR} = 1$ $(1,0)$, $(0,1)$, $(1,1)$
- Red open dot = $\text{OR} = 0$ $(0,0)$
- Dashed line = perceptron decision boundary
- Line separates false from true cases
- OR is linearly separable \rightarrow single-layer perceptron can learn it



Graphical Representation of the AND Function

- Inputs s_1 and s_2 plotted in 2-D plane
- Dashed red line = perceptron decision boundary
- Boundary separates true from false cases
- AND is linearly separable \rightarrow perceptron can learn it

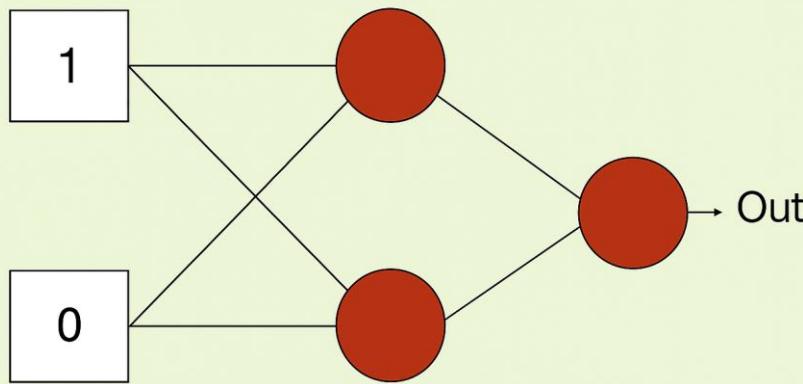


Why XOR Cannot Be Learned by a Single-Layer Perceptron

- ▶ XOR truth table: $(0,1)$ and $(1,0) \rightarrow 1$; $(0,0)$ and $(1,1) \rightarrow 0$
- ▶ Points cannot be separated by a single straight line
- ▶ No linear decision boundary separates true and false cases
- ▶ Therefore: XOR is NOT linearly separable
- ▶ A perceptron can only learn linearly separable functions
- ▶ → XOR requires a multi-layer network (with hidden layer)

The more complex network:

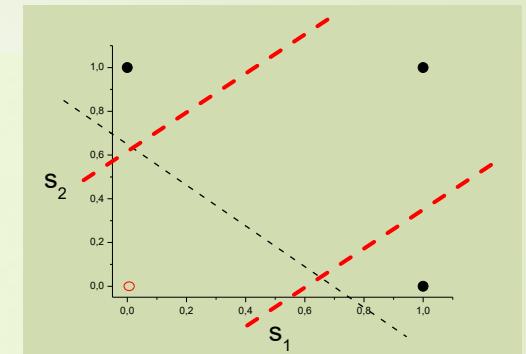
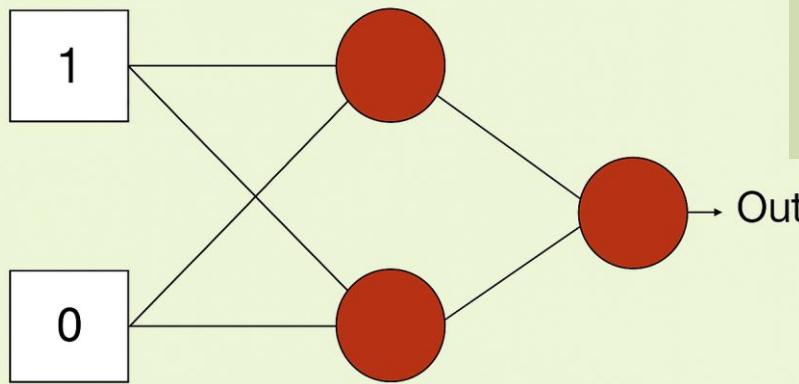
Input layer Hidden layer Output



- ▶ - Hidden Neuron 1 learns one boundary.
- ▶ - Hidden Neuron 2 learns the other boundary.
- ▶ - The output neuron combines the two boundaries into the XOR rule.
- ▶ - This transforms the input so XOR becomes linearly separable at the output layer.
- ▶ - Smallest architecture that solves XOR: 2 input → 2 hidden → 1 output.

The more complex network:

Input layer Hidden layer Output



- ▶ - Hidden Neuron 1 learns one boundary.
- ▶ - Hidden Neuron 2 learns the other boundary.
- ▶ - The output neuron combines the two boundaries into the XOR rule.
- ▶ - This transforms the input so XOR becomes linearly separable at the output layer.
- ▶ - Smallest architecture that solves XOR: 2 input → 2 hidden → 1 output.

A larger network

Input
neurons

Hidden
layer

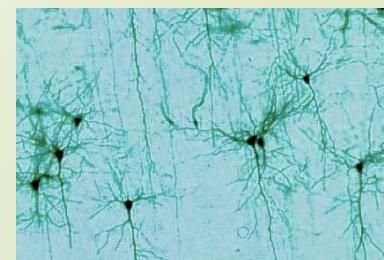
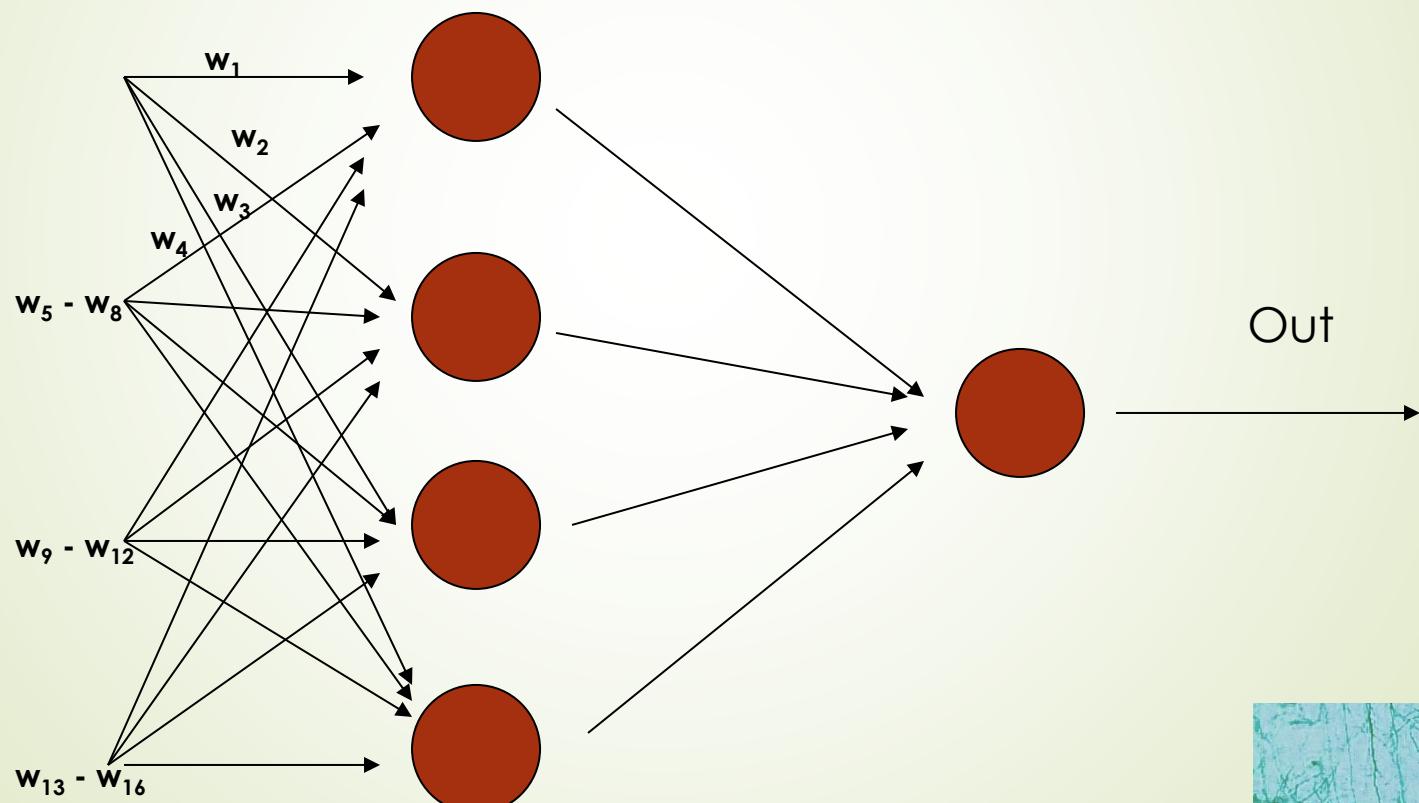
Output
neuron

1

1

0

0

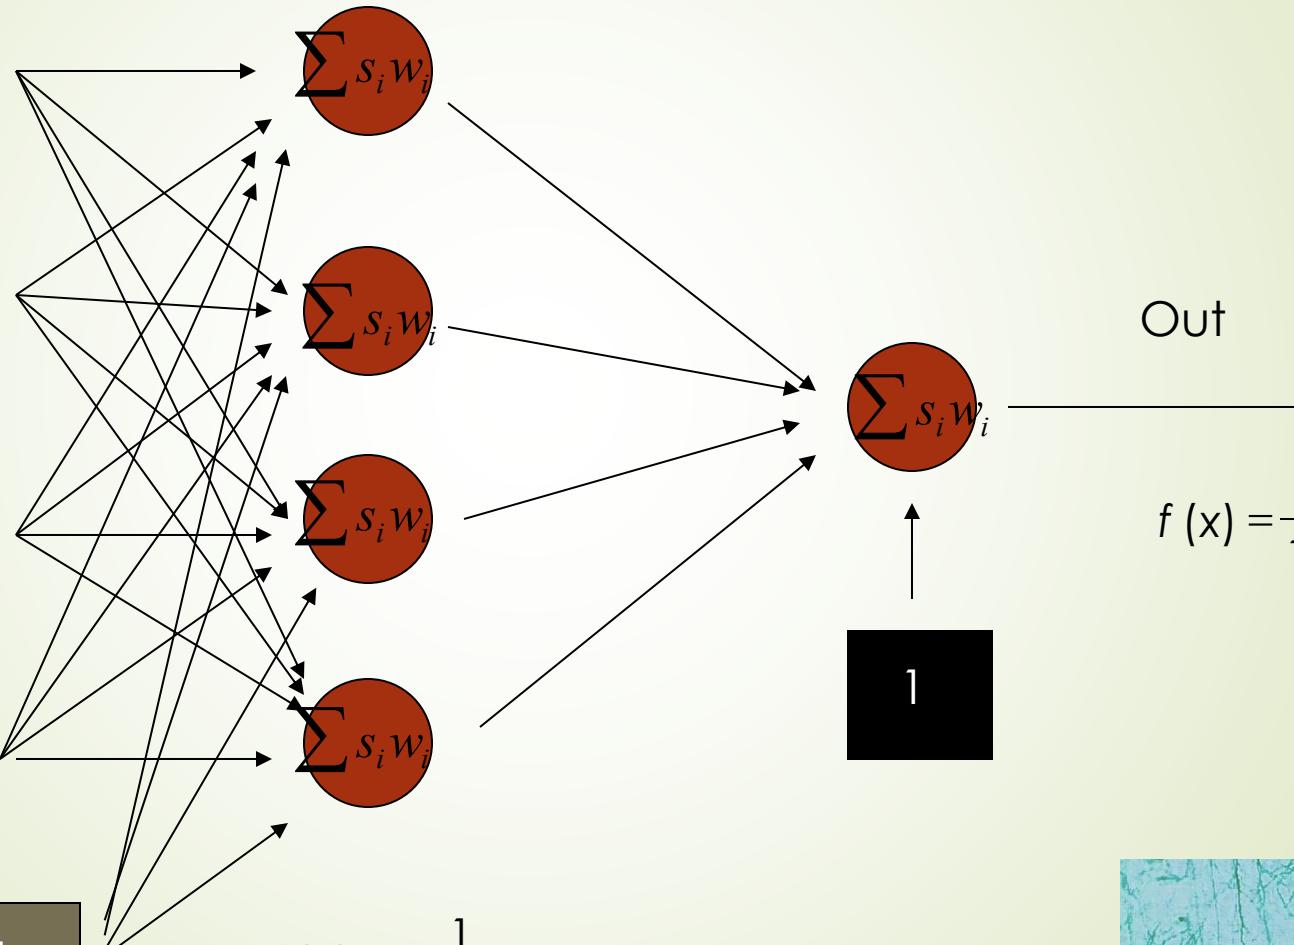
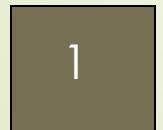


Adding activation functions:

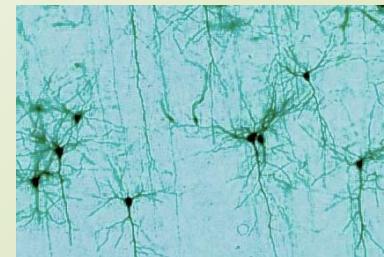
Input
neurons

Hidden
layer

Output
neuron

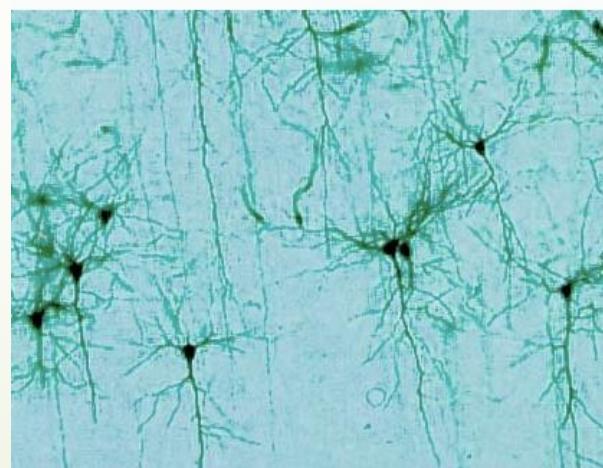
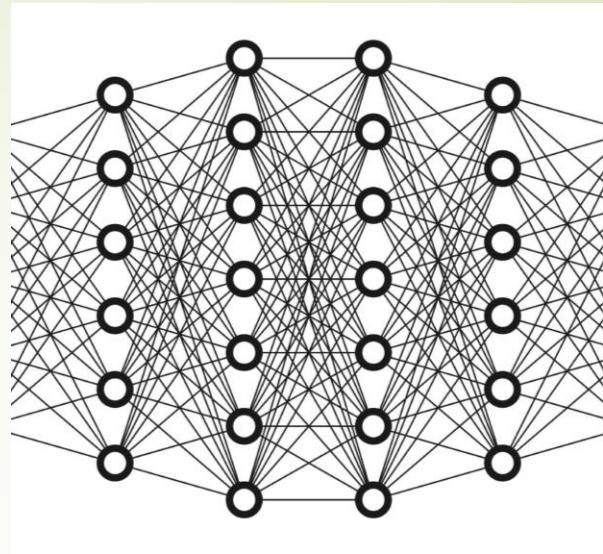


$$f(x) = \frac{1}{1+e^{-x}}$$



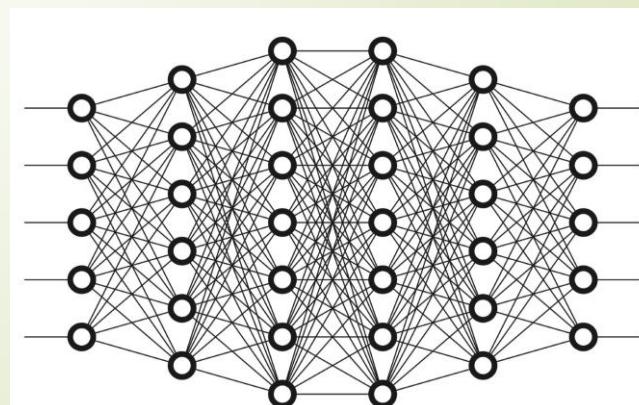
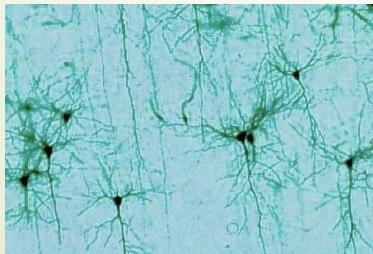
How Hidden Layers Learn Complex Nonlinear Problems

- ▶ Single-layer perceptrons can only learn linear boundaries.
- ▶ Hidden layers introduce nonlinear transformations.
- ▶ Each layer progressively extracts more abstract patterns from input data.
- ▶ Enables solving problems far beyond linear separability.



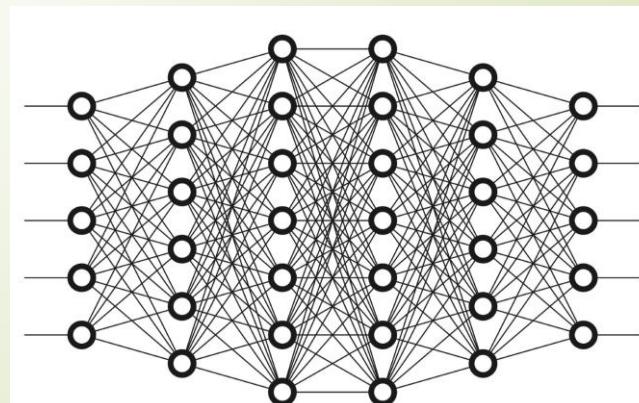
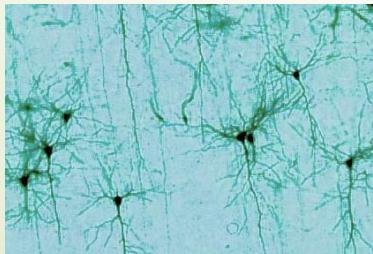
Step 1: Hidden Neurons Create Multiple Linear Boundaries

- ▶ Each hidden neuron computes a weighted sum of inputs.
- ▶ This creates a separate linear decision boundary (line/plane/hyperplane).
- ▶ With H hidden neurons, the network forms H linear partitions.
- ▶ These partitions divide input space into multiple regions.



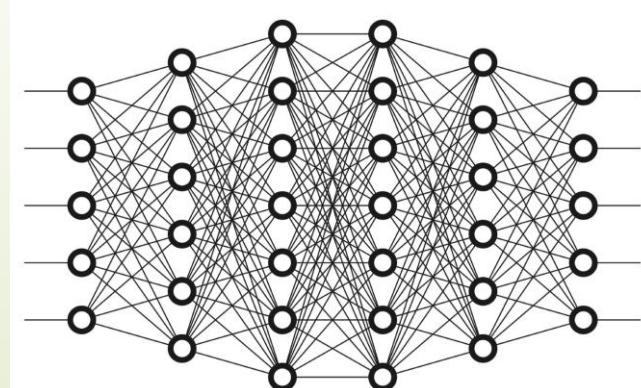
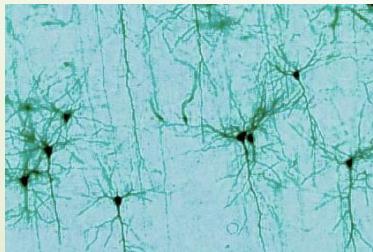
Step 2: Nonlinear Activations Bend and Combine Regions

- ▶ Activations like sigmoid, tanh, and ReLU introduce nonlinearity.
- ▶ They allow the network to warp boundaries into curves.
- ▶ Networks can represent nested, curved, or disjoint regions.
- ▶ The result: flexible and powerful transformation of input space.



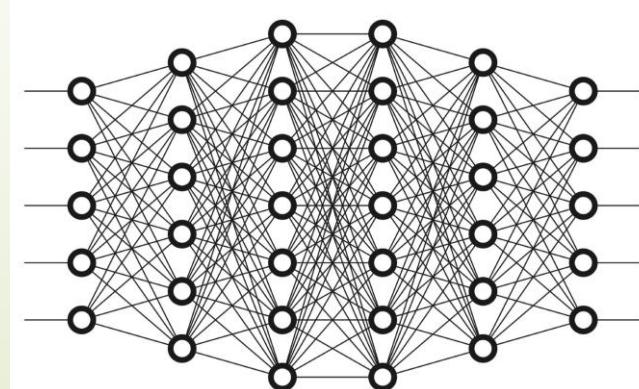
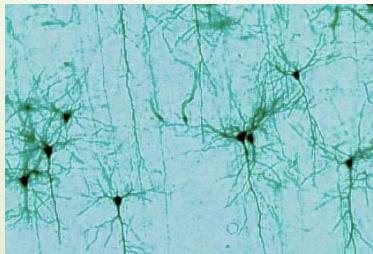
Step 3: Layer-by-Layer Composition Builds Complexity

- ▶ Each deeper layer receives transformed outputs from previous layers.
- ▶ Early layers learn simple features (edges, directions, contrasts).
- ▶ Later layers combine them into complex concepts.
- ▶ This hierarchical composition allows deep networks to approach universal function approximation.



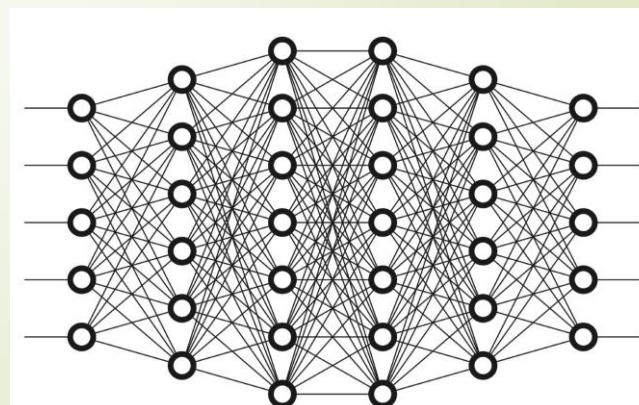
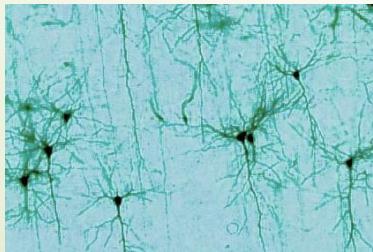
Summary: Why Hidden Layers Enable Advanced Learning

- ▶ Multiple linear boundaries + nonlinear activations = flexible regions.
- ▶ Hidden layers transform data into linearly separable representations.
- ▶ Deep networks can approximate any continuous function.
- ▶ Enables learning of complex patterns in vision, language, biology, and more.



Where is Knowledge Stored in an ANN?

- ▶ Knowledge in an artificial neural network is stored entirely in its weights and biases.
 - ▶ Weights encode the strength of connections between neurons
 - ▶ Biases shift activation thresholds to improve flexibility
 - ▶ Learning = adjusting weights and biases to reduce error
 - ▶ Neurons, activation functions, and architecture do NOT store learned knowledge



What
about
learning?



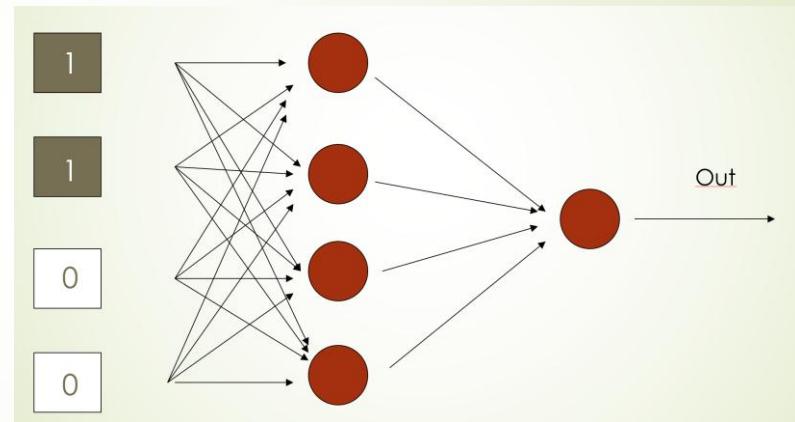


Learning Rules in Neural Networks

- ▶ Hebbian learning (still relevant conceptually): $\Delta w \propto x_i \cdot x_j$ (unsupervised correlation learning).
- ▶ Gradient-based learning (modern standard):
 - Backpropagation: uses chain rule to compute $\partial L / \partial w$ for all weights.
 - Gradient descent update: $\Delta w = -\eta \cdot \partial L / \partial w$.
 - Variants: SGD, Momentum, RMSProp, Adam (most widely used today).
- ▶ Perceptron learning rule: historically important, rarely used today except for teaching.
- ▶ Genetic algorithms: used in specialized cases (neuroevolution)
- ▶ Summary: Modern neural networks are trained almost exclusively with gradient-based optimization + backpropagation.

Feedforward Network

- ▶ Signals flow left → right only
- ▶ Input layer feeds hidden layer; hidden feeds output
- ▶ No loops or feedback connections
- ▶ Used for prediction (classification/regression)
- ▶ Computation: $\text{output} = f(\sum w \cdot x)$
- ▶ Trained by back propagation and gradient descent



Backpropagation Phase

- ▶ Error signal flows backward from output → hidden → input weights
- ▶ Each weight updated based on its contribution to the error
- ▶ Arrows in hidden layer show gradient flow
- ▶ Allows hidden neurons to adjust internal representations
- ▶ Full learning cycle = forward pass + backward pass

Algorithmic Order of Backpropagation

- ▶ 1. Forward Pass: Compute activations layer-by-layer (Input → Hidden → Output).
- ▶ 2. Output Error: Compare prediction to target and compute output-layer error.
- ▶ 3. Backward Pass Step 1: Propagate error from output layer to previous hidden layer.
- ▶ 4. Backward Pass Step 2: Continue propagating error one layer at a time toward the input.
- ▶ 5. Gradient Calculation: Compute gradients for each weight using the layer-specific error.
- ▶ 6. Weight Update: After all gradients are computed, update all weights simultaneously.

Gradients in Backpropagation (Chain Rule Overview)

- ▶ Forward relations (single neuron): $z = \sum w_i x_i + b$, $\hat{y} = f(z)$.
- ▶ Interpretation: z is the weighted input; \hat{y} is the neuron's output after activation f .
- ▶ Goal: Compute how much changing a weight w_i changes the loss $L \rightarrow \partial L / \partial w_i$.
- ▶ Chain rule: $\partial L / \partial w_i = (\partial L / \partial \hat{y}) \cdot (\partial \hat{y} / \partial z) \cdot (\partial z / \partial w_i)$.
 - ▶ $\partial L / \partial \hat{y}$ — Change in loss when the prediction \hat{y} changes (from loss function).
 - ▶ $\partial \hat{y} / \partial z$ — Change in neuron output when input z changes (activation derivative).
 - ▶ $\partial z / \partial w_i$ — Change in weighted sum z when weight w_i changes (equals input x_i).
- ▶ Gradient descent update: $w_i \leftarrow w_i - \eta \cdot \partial L / \partial w_i$.
- ▶ In multi-layer networks: this gradient computation is applied recursively backward through every layer so each neuron adjusts its weights based on its contribution to the final error.

- ▶ Supervised learning: model learns from labelled data ($\text{input} \rightarrow \text{target}$).
- ▶ Training uses loss function to measure prediction error.
- ▶ Backpropagation + gradient descent adjust weights to minimise loss.
- ▶ Unsupervised learning: no explicit labels or target values.
- ▶ Goal: discover structure or patterns directly from the input data.
- ▶ Transition concept: instead of predicting targets, the ANN learns internal structure.

From Supervised to Unsupervised Learning

Backpropagation for Unsupervised Learning

- ▶ Backpropagation still works if we provide a suitable objective.
- ▶ Common approach: self-supervision (model creates its own learning signal).
- ▶ Loss function compares outputs to reconstructed or transformed versions of input data.
- ▶ Gradients computed normally → update weights without external labels.
- ▶ Thus: backprop is a general optimisation tool, not limited to supervised tasks.



Hebbian Learning in Unsupervised Systems

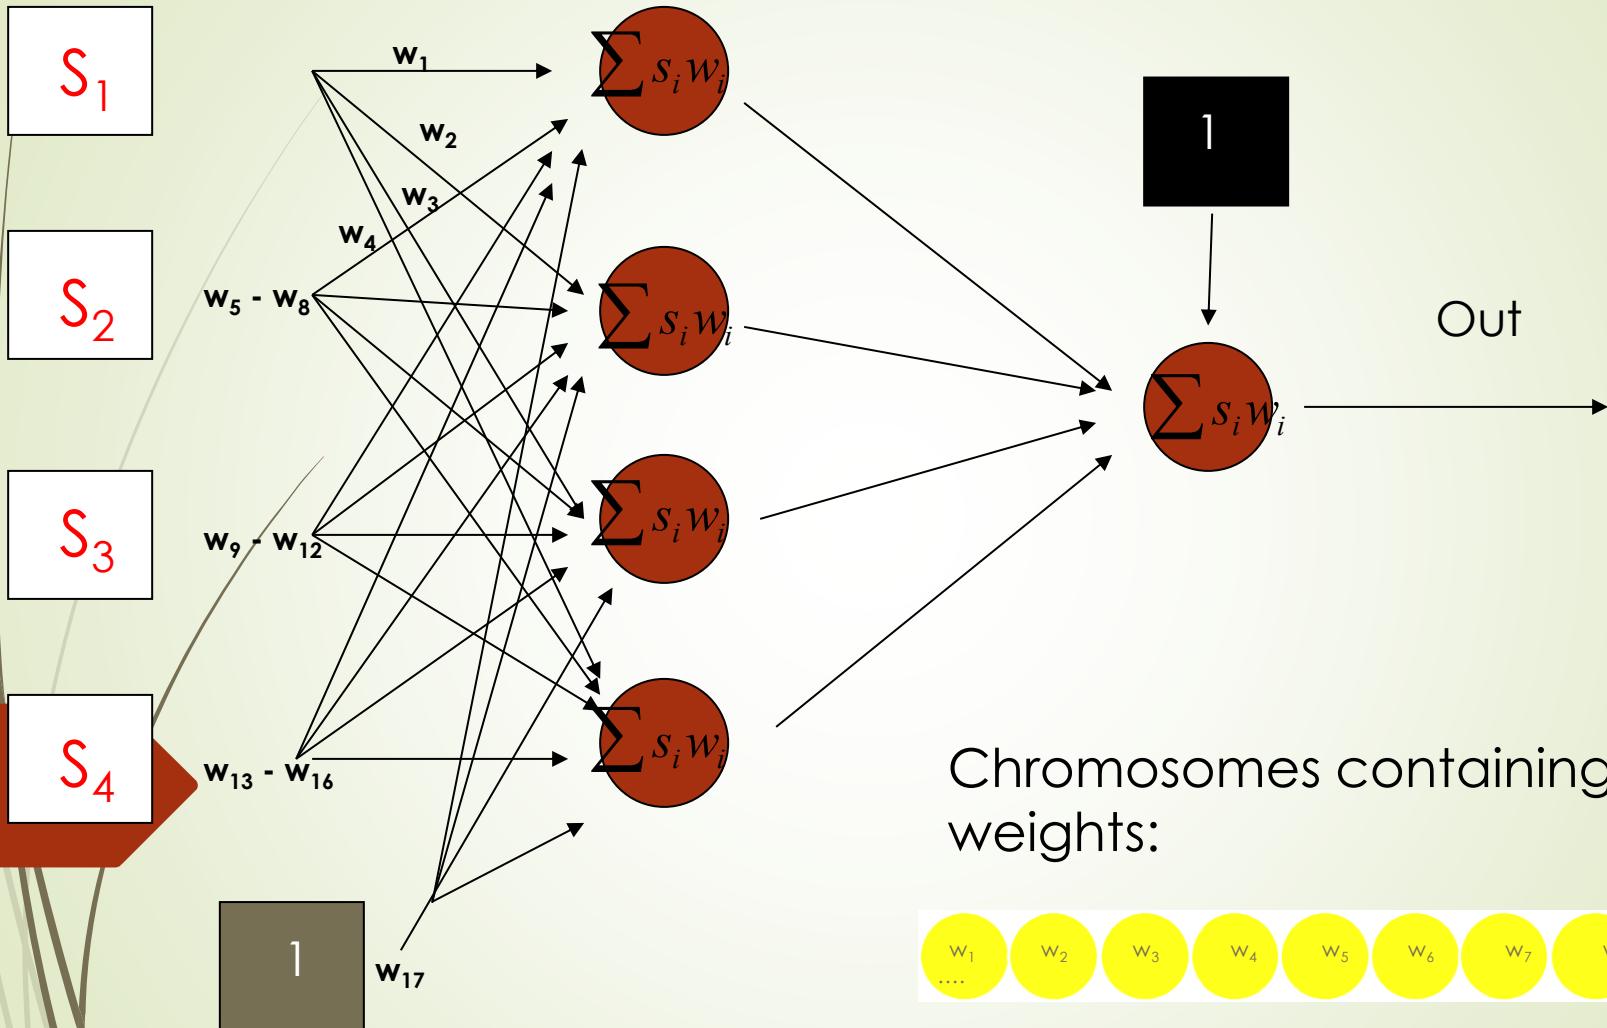
- ▶ Hebbian rule: “cells that fire together, wire together.”
- ▶ Weight update proportional to correlation between neuron activations.
- ▶ $\Delta w_{ij} \propto x_i \cdot x_j$ (or variants such as Oja’s rule).
- ▶ No loss function: learning emerges from co-activation statistics.
- ▶ Useful for feature discovery, clustering, and self-organisation.



Applications of Supervised and Unsupervised Learning

- ▶ Supervised learning:
 - ▶ – Image classification (e.g., crow species).
 - ▶ – Speech recognition, translation.
 - ▶ – Medical diagnosis, regression tasks.
- ▶ Unsupervised learning:
 - ▶ – Clustering (grouping similar individuals).
 - ▶ – Dimensionality reduction (PCA, autoencoders).
 - ▶ – Pattern discovery, anomaly detection.
- ▶ Both approaches complement each other in modern ML systems.

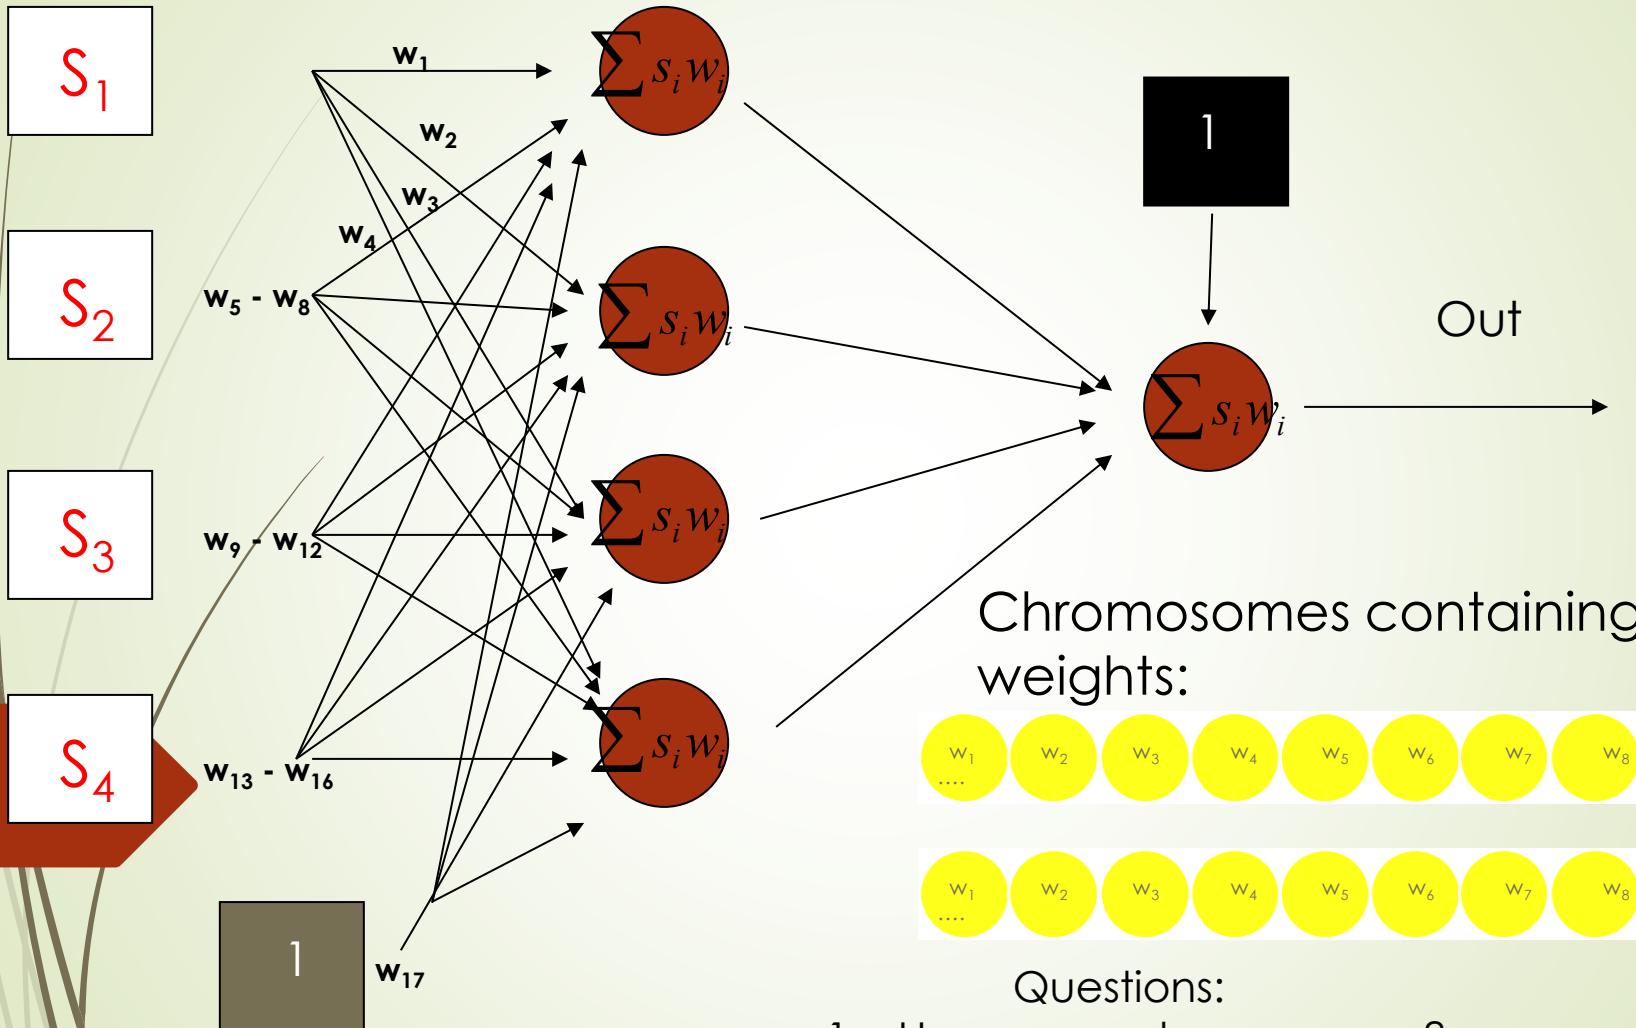
Extra: Neural network optimised by genetic algorithm:



Chromosomes containing weights:



Extra: Neural network optimised by genetic algorithm:



Questions:

1. How many chromosomes?
2. How many genes in each?
3. What type of reproduction?